



SETTING UP A BIOINFORMATICS QC PIPELINE

BRIAN MCCONEGHY

BIOINFORMATICS SPECIALIST

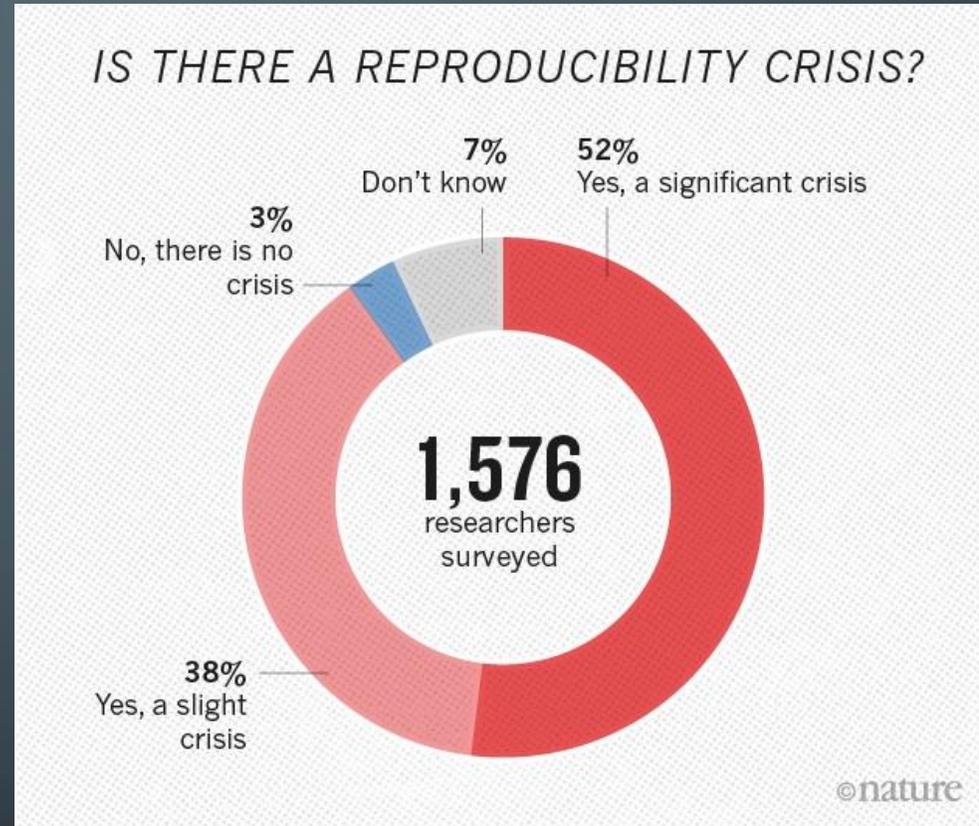
SEQUENCING AND BIOINFORMATICS CONSORTIUM, UBC

OFFICE OF THE VICE-PRESIDENT, RESEARCH & INNOVATION

WESTGRID WEBINAR 2019-11-13

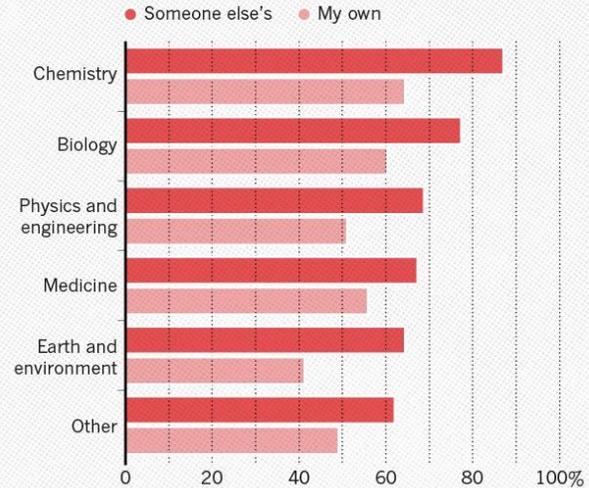
“More than 70% of researchers have tried and failed to reproduce another scientist's experiments, and more than half have failed to reproduce their own experiments.”

Baker, M. (2016). 1,500 scientists lift the lid on reproducibility. *Nature*, 533(7604), 452-454. doi:10.1038/533452a



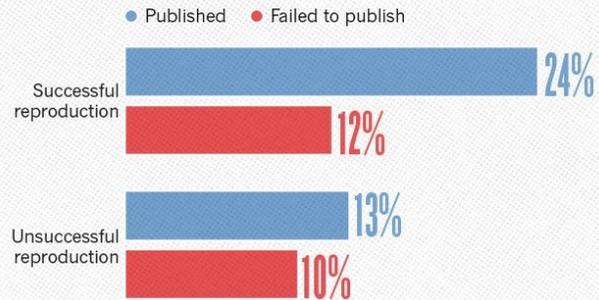
HAVE YOU FAILED TO REPRODUCE AN EXPERIMENT?

Most scientists have experienced failure to reproduce results.



HAVE YOU EVER TRIED TO PUBLISH A REPRODUCTION ATTEMPT?

Although only a small proportion of respondents tried to publish replication attempts, many had their papers accepted.



Number of respondents from each discipline:
Biology **703**, Chemistry **106**, Earth and environmental **95**,
Medicine **203**, Physics and engineering **236**, Other **233** 



Background

Pipelining Tools

Writing the Pipeline

Metrics to Track

Implementation

Conclusions



Background

Pipelining Tools

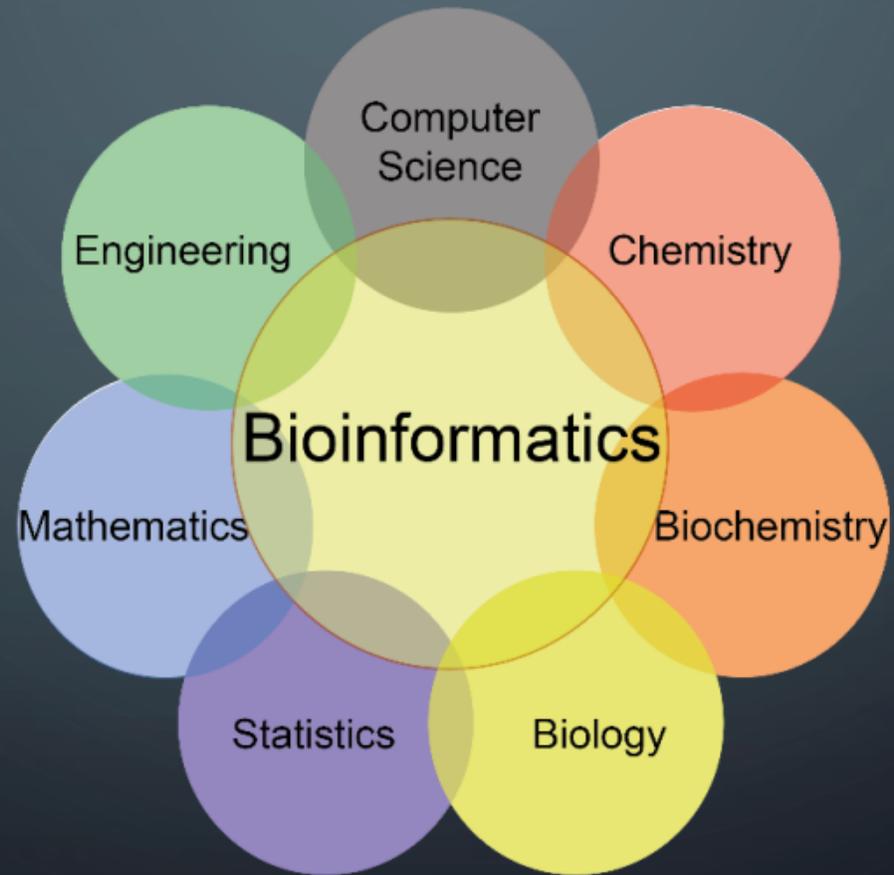
Writing the Pipeline

Metrics to Track

Implementation

Conclusions

WHAT IS BIOINFORMATICS?



NEXT GENERATION SEQUENCING

Sample (input) QC

Sample Preparation

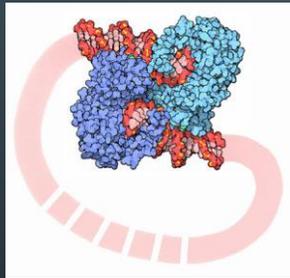
Sample (library) QC

Sequencing

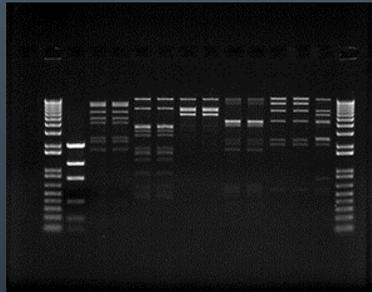
Data!



<https://www.nextadvance.com/>



<https://pdb101.rcsb.org/motm/84>



www.thermofisher.com



www.illumina.com



<https://www.makeuseof.com/tag/best-linux-server-operating-systems/>

NEXT GENERATION SEQUENCING

Sample (input) QC

Sample Preparation

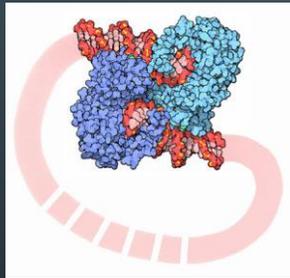
Sample (library) QC

Sequencing

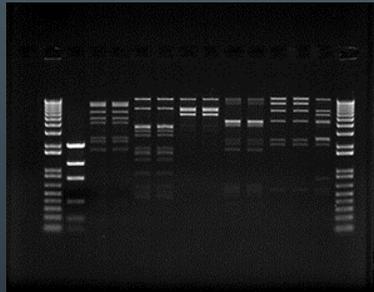
Data!



<https://www.nextadvance.com/>



<https://pdb101.rcsb.org/motm/84>



www.thermofisher.com



www.illumina.com



<https://www.makeuseof.com/tag/best-linux-server-operating-systems/>

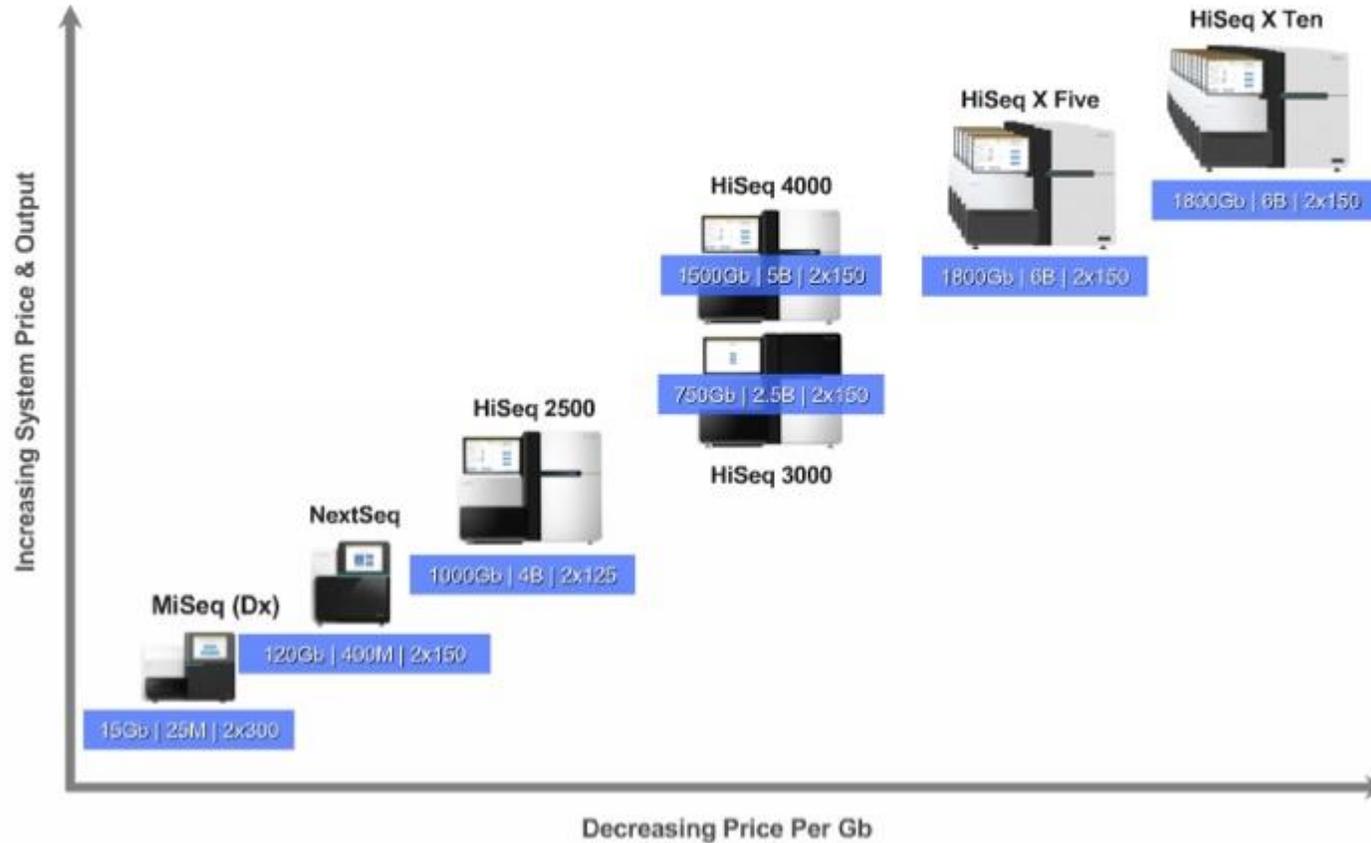
WHAT IS NEXT GEN SEQUENCING DATA, EXACTLY?

- High-throughput sequencing technology
 - Generates **millions** of 'reads'
 - Reads are just strings of G's, A's, T's, and C's (with associated quality values)

```
@NB999999:999:ZZZZZZZ9:1:11101:16570:1094 1:N:0:1  
AAAGCNGCTGAATTGTTTCGCGTTTACCTTGCGTGTACGCGCAGGAAACACT  
+  
AAAAA#A6EA66EEEEEEE/E//EAE/E//EEEEEEEEAAEEEEEEEEEEE
```

phiX 174 control DNA

Sequencing Power For Every Scale.





Background

Pipelining Tools

Writing the Pipeline

Metrics to Track

Implementation

Conclusions



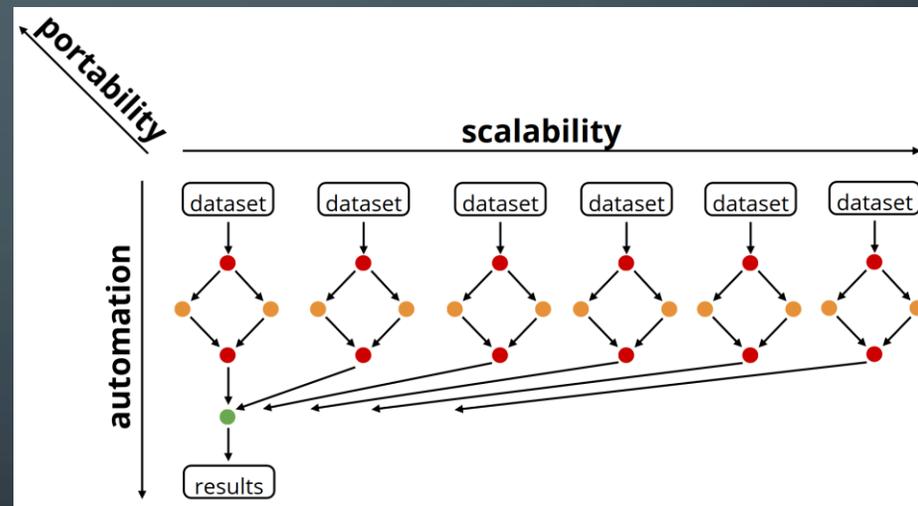
WHAT AND WHY

- Workflow management system
- Reproducible and scalable data analysis
- Rules, inputs, and outputs



NEEDS

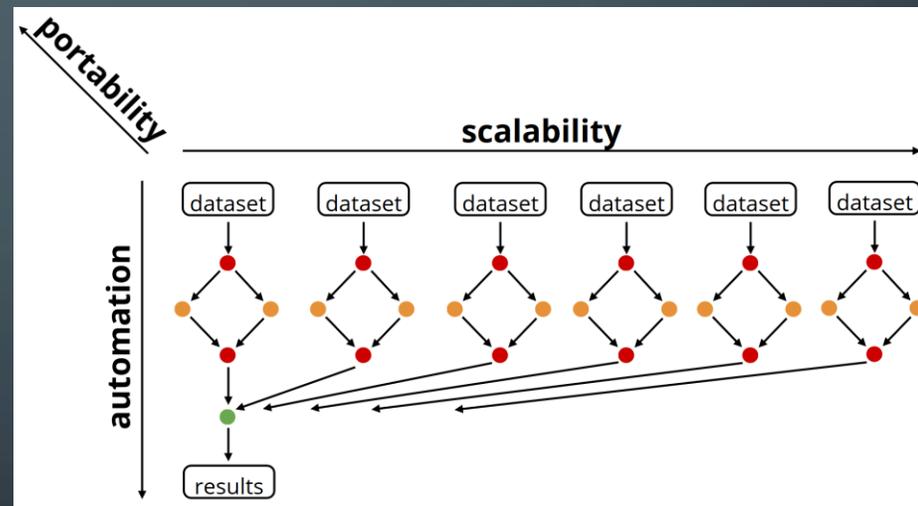
- Reproducible
- Scalable
- Efficient (Parallelizable)
- Portable
- Automated



<https://slides.com/johanneskoester/snakemake-short#/3>

WANTS

- Ease of development
- Unix-compatible
- FREE



<https://slides.com/johanneskoester/snakemake-short#/3>

COMPARISON

- Galaxy
- Ruffus
- Snakemake

COMPARISON

- Galaxy
- Ruffus
- **Snakemake**



Background

Pipelining Tools

Writing the Pipeline

Metrics to Track

Implementation

Conclusions



RESOURCES - HARDWARE

- Cedar - Compute Canada
 - 58,416 Cores
 - 306,306 GB of RAM
 - 10TB scratch space



<https://medium.com/monplan/how-we-automated-deployments-and-testing-with-bitbucket-pipelines-bb478c12c55f>

RESOURCES - PEOPLE

- Advanced Research Computing (ARC)
 - Jamie Rosner
 - Venkat Mahadevan
- VP Research & Innovation (VPRI)
 - Dr. Helen Burt

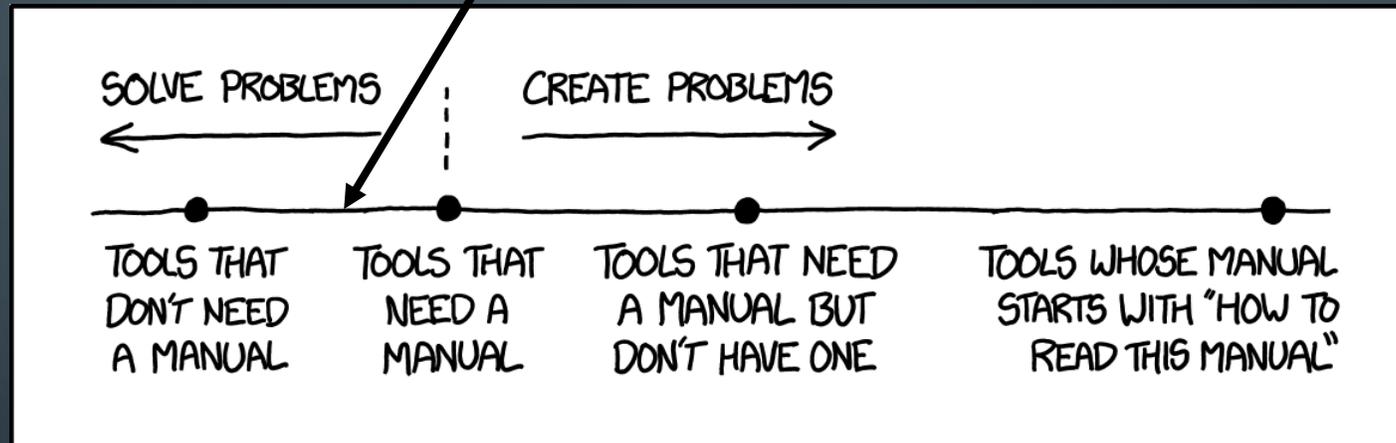


<https://medium.com/monplan/how-we-automated-deployments-and-testing-with-bitbucket-pipelines-bb478c12c55f>

SNAKEMAKE

- Decompose workflow into rules
- Rules define how to obtain output files from input files
- Snakemake infers dependencies and execution order

SLAKEMAKE



<https://xkcd.com/1343/>

And now
for something
completely TECHNICAL ...



SIMPLICITY!

```
rule sort:  
  input:  
    "path/to/dataset.txt"  
  output:  
    "dataset.sorted.txt"  
  shell:  
    "sort {input} > {output}"
```

GENERALIZE RULES WITH NAMED **WILDCARDS**

```
rule sort:  
  input:  
    "path/to/{dataset}.txt"  
  output:  
    "{dataset}.sorted.txt"  
  shell:  
    "sort {input} > {output}"
```

SPECIFY MULTIPLE INPUTS (AND OUTPUTS) REFER BY **INDEX**

```
rule sort_and_annotate:  
  input:  
    "path/to/{dataset}.txt",  
    "path/to/annotation.txt"  
  output:  
    "{dataset}.sorted.txt"  
  shell:  
    "paste <(sort {input[0]}) {input[1]} > {output}"
```

CAN SPECIFY MULTIPLE INPUTS (AND OUTPUTS), AND REFER BY **NAME**

```
rule sort_and_annotate:  
  input:  
    a="path/to/{dataset}.txt",  
    b="path/to/annotation.txt"  
  output:  
    "{dataset}.sorted.txt"  
  shell:  
    "paste <(sort {input.a}) {input.b} > {output}"
```

USE PYTHON WITHIN RULES

```
rule sort:
    input:
        a="path/to/{dataset}.txt"
    output:
        b="{dataset}.sorted.txt"
    run:
        with open(output.b, "w") as out:
            for l in sorted(open(input.a)):
                print(l, file=out)
```

A REAL RULE

- Used in DNA QC pipeline

```
rule bwa_mem_map_reads:
  input:
    get_trimmed_reads
  output:
    temp('mapped/{sample}-{unit}.sorted.bam')
  log:
    'logs/bwa_mem/{sample}-{unit}.log'
  params:
    index = get_genome_index,
    rg = get_read_group_bwa
  threads: 46
  shell:
    '(bwa mem -t {threads} {params.rg} {params.index} {input} | '
    'samtools sort -T $SLURM_TMPDIR/ -o {output} -) 2> {log}'
```

JOB EXECUTION

- A job only executes if:
 1. output file is the target requested and does not exist
 2. output file needed by another executed job (i.e. is an input to another job) and does not exist
 3. input file is newer than the output file
 4. input file will be updated by other job
 5. execution is forced

CLUSTER EXECUTION

- Can set up pipeline **profiles**
- Execute DAG by way of cluster job submission
- Configuration file
 - Max jobs at a time
 - CPUs
 - MEM
 - General (per profile) or granular (per rule)



Background

Pipelining Tools

Metrics to Track

Writing the Pipeline

Implementation

Conclusions



METRICS TO TRACK

- Adapter trimming
- Duplicate Rate
- % Aligned (for genomes we can map to)
- Insert size
- Coverage
- Error rate
- GC Content
- For RNA, specifically:
 - Strand specificity (% correct strand)
 - 5'-3' bias
 - % rRNA
 - Intron-exon ratio



Background

Pipelining Tools

Writing the Pipeline

Metrics to Track

Implementation

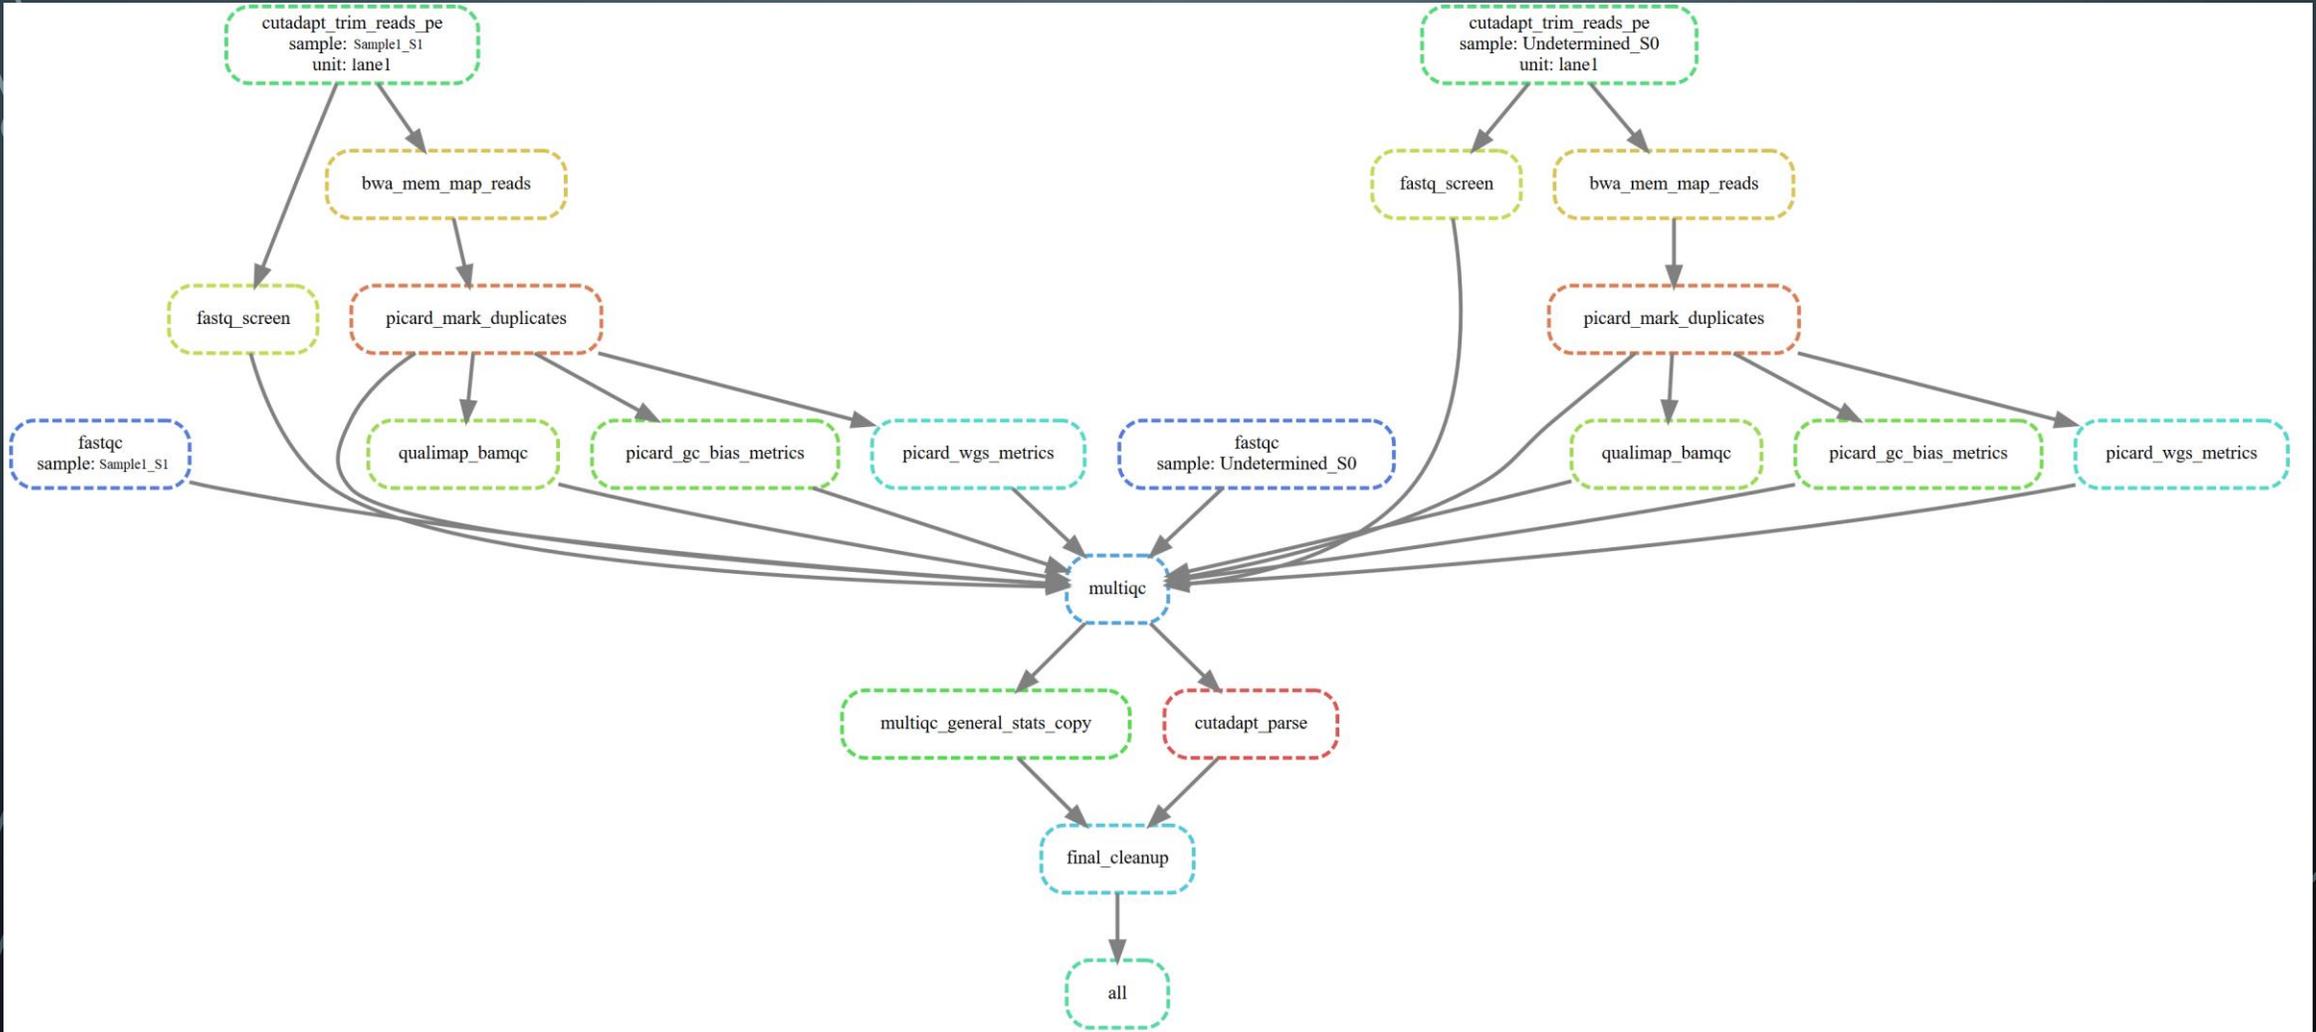
Conclusions



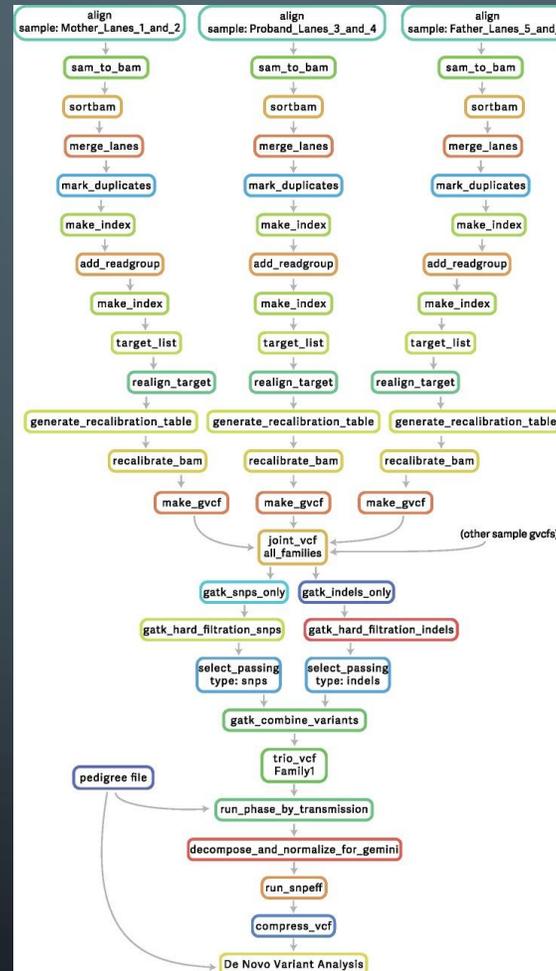
IMPLEMENTATION

- Conda environment
- Version controlled - **GitHub**
- SBC has 3 QC pipelines (combined into 1, dynamically determined)
 - Paired-end DNA QC
 - Single-end DNA QC
 - Paired-end RNA QC

DNA QC WORKFLOW – 2 SAMPLES



WORKFLOWS CAN BE COMPLEX





Background

Pipelining Tools

Writing the Pipeline

Metrics to Track

Implementation

Conclusions

CONCLUSIONS

- Snakemake satisfied all needs of the SBC and is simple to work with
- The complex metrics the SBC is most interested in are being tracked, in an automated fashion
- Implementation allows for reproducible, scalable, flexible, trackable QC

THANK YOU



THE UNIVERSITY OF BRITISH COLUMBIA

Sequencing + Bioinformatics Consortium

Link to GitHub with workshop instructions:

<https://bit.ly/2Xf4HN6>