

# Using Singularity Containers in Virtual Machines and HPC

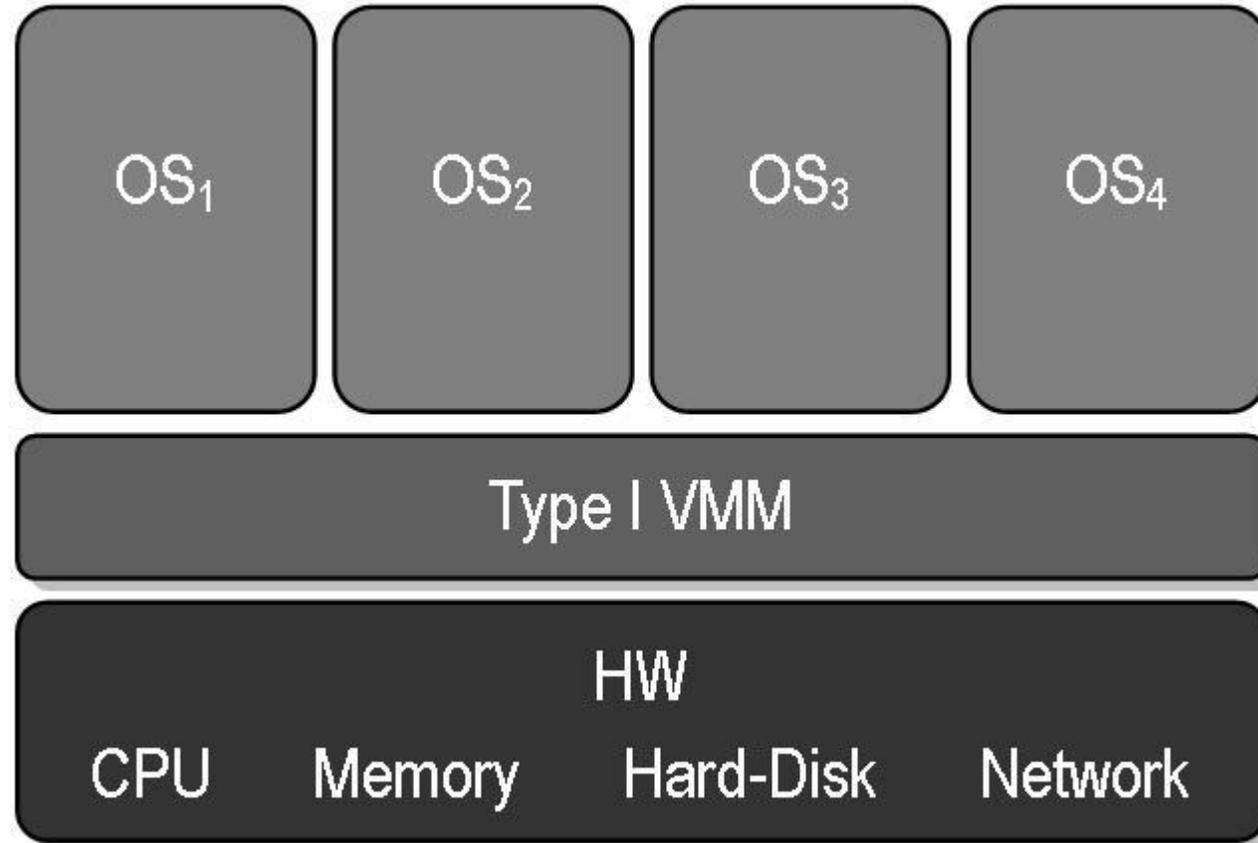
Venkat Mahadevan

UBC Advanced Research Computing (ARC)

5/26/2021



# Virtualization

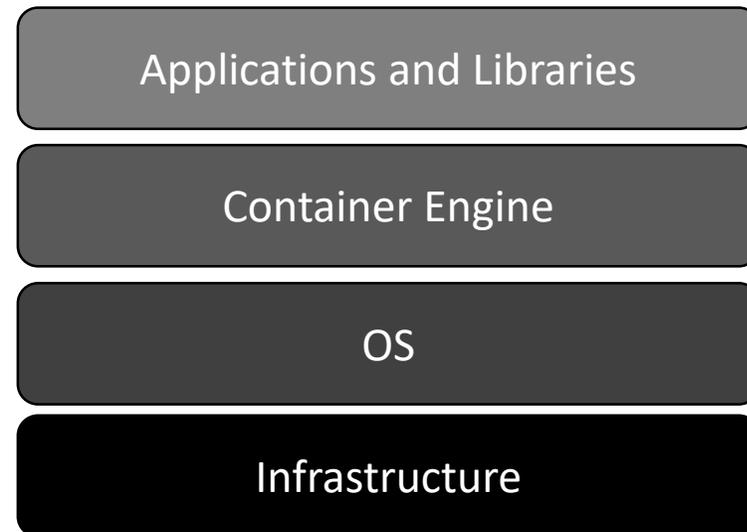


# Containers



# Containers

- Containers provide an additional layer of abstraction over virtualization.



# Containers

- However, containers do not require virtual machines and can happily run on bare-metal.
- Lightweight compared to VMs in terms of memory and storage requirements.
- Increases portability of applications between datacenters, private clouds, and public clouds.



# Containers

- Can provide a very customized environment specific to research use cases.
- Commercial software that requires old, obsolete libraries or operating system versions.
- Bottom line: use containers where it makes sense to your workflow.



# Why Use Containers

- Reproducibility.
- Portability.
- Isolation.
- Avoiding complexity aka dependency hell.

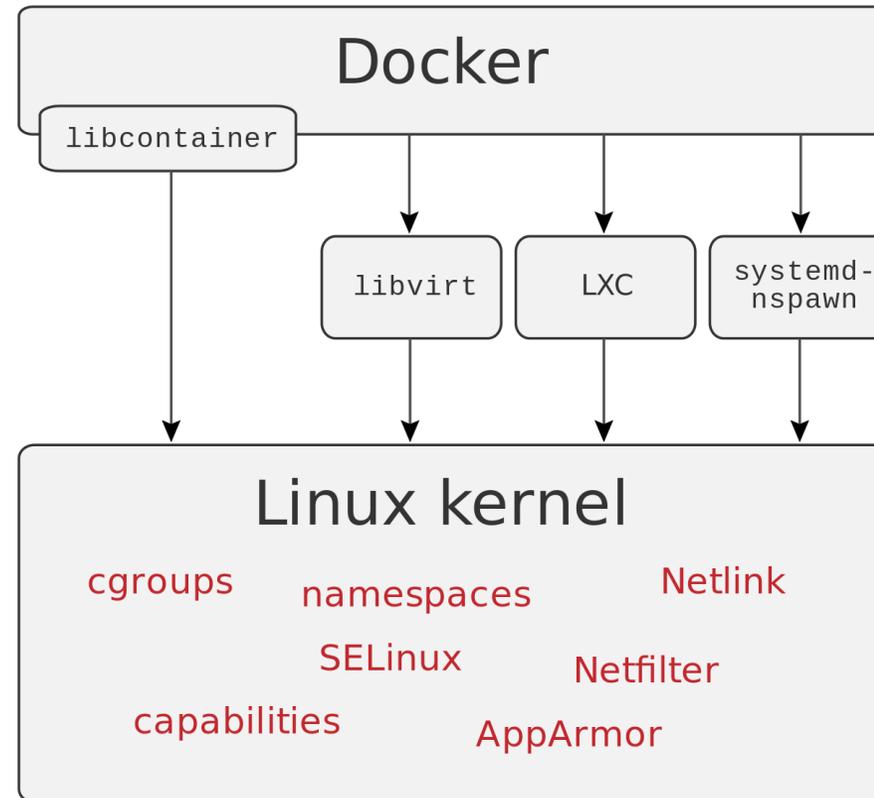


# Docker

- The advent of Docker resulted in widespread adoption of container technology.
- Uses cgroups, kernel namespaces, OverlayFS (union filesystem) to allow containers to run in a singular Linux instance.
- Very mature technology with ~7 million DockerHub users.



# Docker



# Singularity

- However, Docker is not a perfect fit in the advanced research computing space.
- The Docker daemon needs to run in the background with elevated privileges on every node (in the datacenter, private/public Cloud VM, etc.) that hosts containers.
- This is a security concern which can be mitigated but limits some features.



# Singularity

- Docker also does not provide close integration with the standard HPC software stack such as MPI as well as schedulers and resource managers such as Slurm & PBS Pro.
- Container orchestration technologies such as Kubernetes can replicate some of the aforementioned functionality but the existing HPC stack is well entrenched and robust.



# Singularity features

- Was developed at Lawrence Berkeley National Lab specifically for HPC workloads.
- No daemon process required; an executable is provided.
- The user permissions are maintained both in and outside the container; mitigates privilege escalation concerns.



# Singularity features

- Support for both NVIDIA and AMD GPUs: pass the `--nv` or `--rocm` flag to the singularity command line.
- The host and container must have compatible versions of the GPU drivers and libraries installed.
- Support for MPI via a couple of different models: the hybrid model and the bind model.



# Singularity features

- MPI via the Hybrid model:
  - Both the host and the container must have compatible versions of MPI libraries and runtime installed.
  - Simplifies the running of MPI applications (just run the singularity command after mpirun or srun if using Slurm).
- MPI via the Bind model:
  - Mount the MPI implementation of the host in the container.
- <https://sylabs.io/guides/3.7/user-guide/mpi.html>

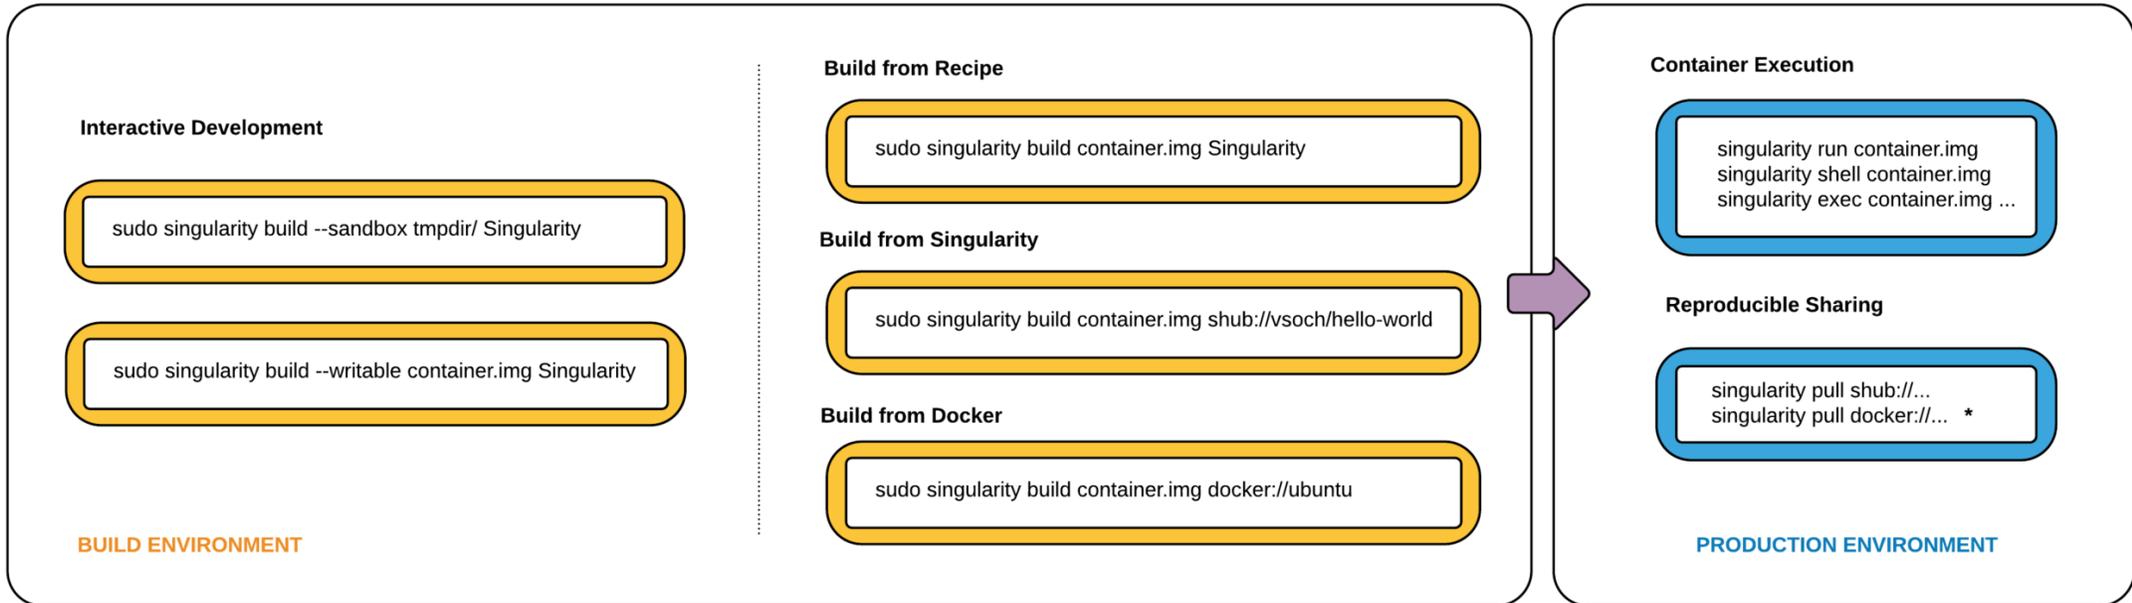


# What about performance?

- On bare-metal, there is almost no performance gap between running in a Singularity container vs. directly on the host system.
- Papers benchmarking the performance of Singularity with various applications have found near native performance (within 1-2% performance) on bare-metal HPC systems:
  - <https://ieeexplore.ieee.org/abstract/document/8855563>
  - <https://people.csail.mit.edu/dpaz/img/PDF/a065-LePaz.pdf>



# Typical workflow



\* Docker construction from layers not guaranteed to replicate between pulls



# DockerHub



## tensorflow/tensorflow ☆

↓ Pulls **50M+**

By [tensorflow](#) • Updated 44 minutes ago

Official Docker images for the machine learning framework TensorFlow (<http://www.tensorflow.org>)

Container

Overview

**Tags**

Filter Tags

Sort by Newest

TAG

[nightly-gpu-jupyter](#)

Last pushed 44 minutes ago by [tensorflowpackages](#)

docker pull tensorflow/tensorflow:nightly-gpu-jupyter

DIGEST

[33f893623392](#)

[33f893623392](#)

OS/ARCH

linux/amd64

linux/amd64

COMPRESSED SIZE ⓘ

2.65 GB

2.65 GB



# Virtual Machines and HPC

- There is always going to be a performance impact of virtualization vs bare-metal.
- However, running containers on virtual machines and whole HPC clusters on VMs and Clouds is gaining popularity.
- Performance impacts can be mitigated by tuning VMs and underlying hardware (e.g. AWS Elastic Network/Fabric Adapter).



# Deep Learning with Horovod

- Horovod is a distributed deep learning training framework for TensorFlow, Keras, PyTorch, and Apache MXNet.
- <https://github.com/horovod/horovod/>
- The stated goal is to take single-GPU training and scale it across many GPUs in parallel.



# Distributed training example

- Build singularity container using DockerHub:

```
singularity pull horovod.sif docker://horovod/horovod:sha-40cbc9b
```

- Run on a single GPU:

```
singularity exec --nv horovod.sif horovodrun -np 1 -H localhost:1  
python ~/pytorch_imagenet_resnet50.py --train-dir ~/tiny-imagenet-  
200/train --val-dir ~/tiny-imagenet-200/val
```



# Distributed training example

- Will run the example on a virtual HPC on AWS Parallel Cluster.
- 8 GPU nodes with V100 running on AWS EC2 P3 instances.
- Slurm scheduler with Open MPI.



The screenshot shows the AWS Parallel Cluster console for a cluster named 'testpackage.ubc-hpc...'. The cluster is in a 'V' (Verified) state and has 22 SSH connections available, managed by the 'ubuntu' user. The cluster's RPID is 'AVRCDEV:vm.pub'.

The console displays two node types:

- MASTER NODE - SCHED...:** P3.2XLARGE instance type. It has 61 GiB RAM, 8 vCPUs, and 1 GPU. The operating system is Ubuntu Server 18.04 LTS, and the cluster version is v2.6.1.
- COMPUTE NODES:** P3.2XLARGE instance type. It has 61 GiB RAM, 8 vCPUs, and 1 GPU. The operating system is Ubuntu Server 18.04 LTS. A green badge indicates that 8 of 10 compute nodes are currently running.

# Distributed training example

```
2021-05-26 14:16:22.116100: I tensorflow/stream_executor/platform/default/dso_loader.cc:53] Successfully opened dynamic library libcudart.so.11.0
Train Epoch #7: 1% | 3/391 [00:02<04:36, 1.40it/s, loss=3.4, accuracy=22.9]
Train Epoch #7: 2% || 6/391 [00:03<02:58, 2.16it/s, loss=3.5, accuracy=21.4]
Train Epoch #7: 2% || 9/391 [00:05<02:30, 2.53it/s, loss=3.46, accuracy=22]
Train Epoch #7: 3% ||| 12/391 [00:06<02:20, 2.69it/s, loss=3.46, accuracy=23.1]
Train Epoch #7: 4% ||| 16/391 [00:07<02:15, 2.76it/s, loss=3.45, accuracy=23.1]
Train Epoch #7: 5% ||| 19/391 [00:08<02:17, 2.71it/s, loss=3.45, accuracy=23.4]
Train Epoch #7: 6% ||| 22/391 [00:09<02:18, 2.66it/s, loss=3.43, accuracy=23.7]
Train Epoch #7: 6% ||| 25/391 [00:11<02:16, 2.68it/s, loss=3.43, accuracy=23.6]
Train Epoch #7: 7% ||| 28/391 [00:12<02:17, 2.64it/s, loss=3.44, accuracy=23.4]
Train Epoch #7: 8% ||| 31/391 [00:13<02:14, 2.67it/s, loss=3.44, accuracy=23.6]
Train Epoch #7: 9% ||| 35/391 [00:14<02:10, 2.74it/s, loss=3.43, accuracy=23.7]
Train Epoch #7: 10% ||| 38/391 [00:15<02:10, 2.71it/s, loss=3.41, accuracy=24.1]
Train Epoch #7: 10% ||| 41/391 [00:16<02:14, 2.60it/s, loss=3.4, accuracy=24.2]
Train Epoch #7: 11% ||| 44/391 [00:18<02:21, 2.45it/s, loss=3.4, accuracy=24.3]
Train Epoch #7: 12% ||| 47/391 [00:19<02:12, 2.59it/s, loss=3.41, accuracy=24.2]
Train Epoch #7: 13% ||| 50/391 [00:20<02:17, 2.48it/s, loss=3.41, accuracy=24.2]
Train Epoch #7: 14% ||| 54/391 [00:21<02:08, 2.62it/s, loss=3.4, accuracy=24.3]
Train Epoch #7: 15% ||| 57/391 [00:23<02:10, 2.56it/s, loss=3.41, accuracy=24.2]
Train Epoch #7: 15% ||| 60/391 [00:24<02:09, 2.56it/s, loss=3.41, accuracy=24.3]
Train Epoch #7: 16% ||| 63/391 [00:25<02:08, 2.55it/s, loss=3.41, accuracy=24.2]
Train Epoch #7: 17% ||| 66/391 [00:26<02:04, 2.60it/s, loss=3.41, accuracy=24.2]
Train Epoch #7: 18% ||| 69/391 [00:28<02:03, 2.60it/s, loss=3.41, accuracy=24.3]
Train Epoch #7: 18% ||| 72/391 [00:29<01:57, 2.72it/s, loss=3.4, accuracy=24.3]
Train Epoch #7: 19% ||| 75/391 [00:30<01:53, 2.78it/s, loss=3.4, accuracy=24.4]
Train Epoch #7: 20% ||| 78/391 [00:31<01:46, 2.95it/s, loss=3.39, accuracy=24.3]
Train Epoch #7: 21% ||| 82/391 [00:32<01:50, 2.81it/s, loss=3.4, accuracy=24.3]
:51, 2.78it/s, loss=3.4, accuracy=24.3]
ubuntu@ip-10-255-2-228:~$
```



TTK Demo on HPC

# Topology ToolKit

Efficient, generic and easy  
Topological data analysis and visualization



# References

- Singularity on Compute Canada HPC:  
<https://docs.computecanada.ca/wiki/Singularity>
- Sylabs.io Singularity Official User Guide:  
<https://sylabs.io/guides/3.7/user-guide/>
- Westgrid YouTube channel:  
<https://www.youtube.com/channel/UCfgds4Qf7VFOv4ORRvFFmhw>



# Attributions

- [https://en.wikipedia.org/wiki/Intermodal\\_container#/media/File:Line3174 - Shipping Containers at the terminal at Port Elizabeth, New Jersey - NOAA.jpg](https://en.wikipedia.org/wiki/Intermodal_container#/media/File:Line3174_-_Shipping_Containers_at_the_terminal_at_Port_Elizabeth,_New_Jersey_-_NOAA.jpg)
- <https://commons.wikimedia.org/wiki/File:VMM-Type1.JPG>
- [https://en.wikipedia.org/wiki/Docker\\_\(software\)#/media/File:Docke-linux-interfaces.svg](https://en.wikipedia.org/wiki/Docker_(software)#/media/File:Docke-linux-interfaces.svg)
- <http://singularity.lbl.gov/assets/img/diagram/singularity-2.4-flow.png>

