

Singularity containers for Bioinformatics workflows

Dr. Tannistha Nandi

Data Scientist

Research Computing Services, University of Calgary

Email: tannistha.nandi@compute.canada.ca

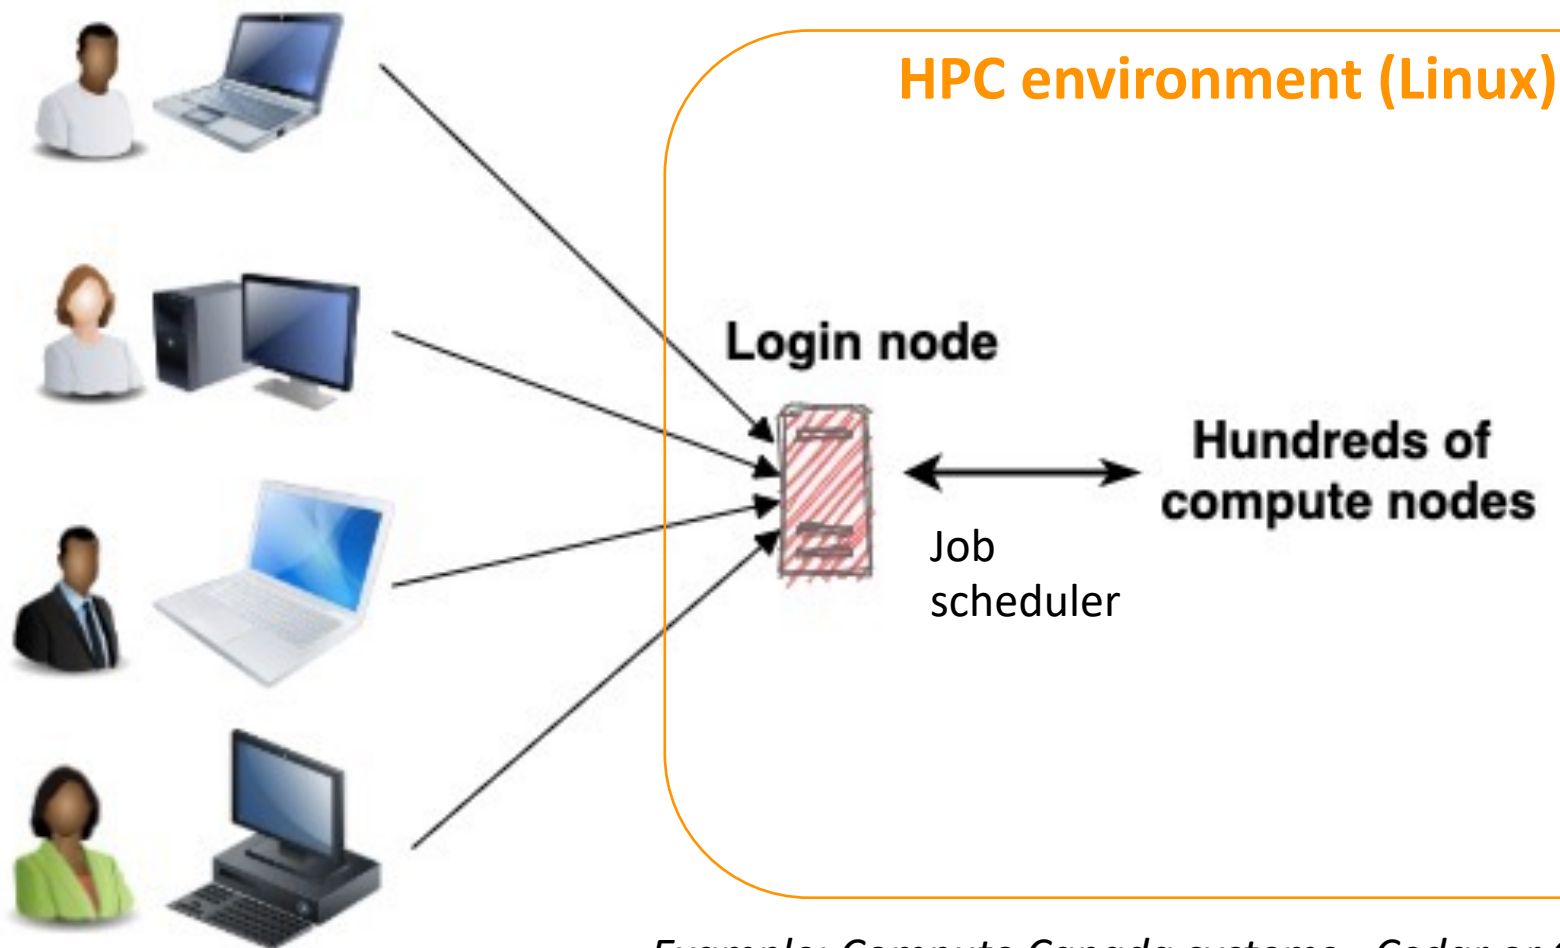


Goals for today

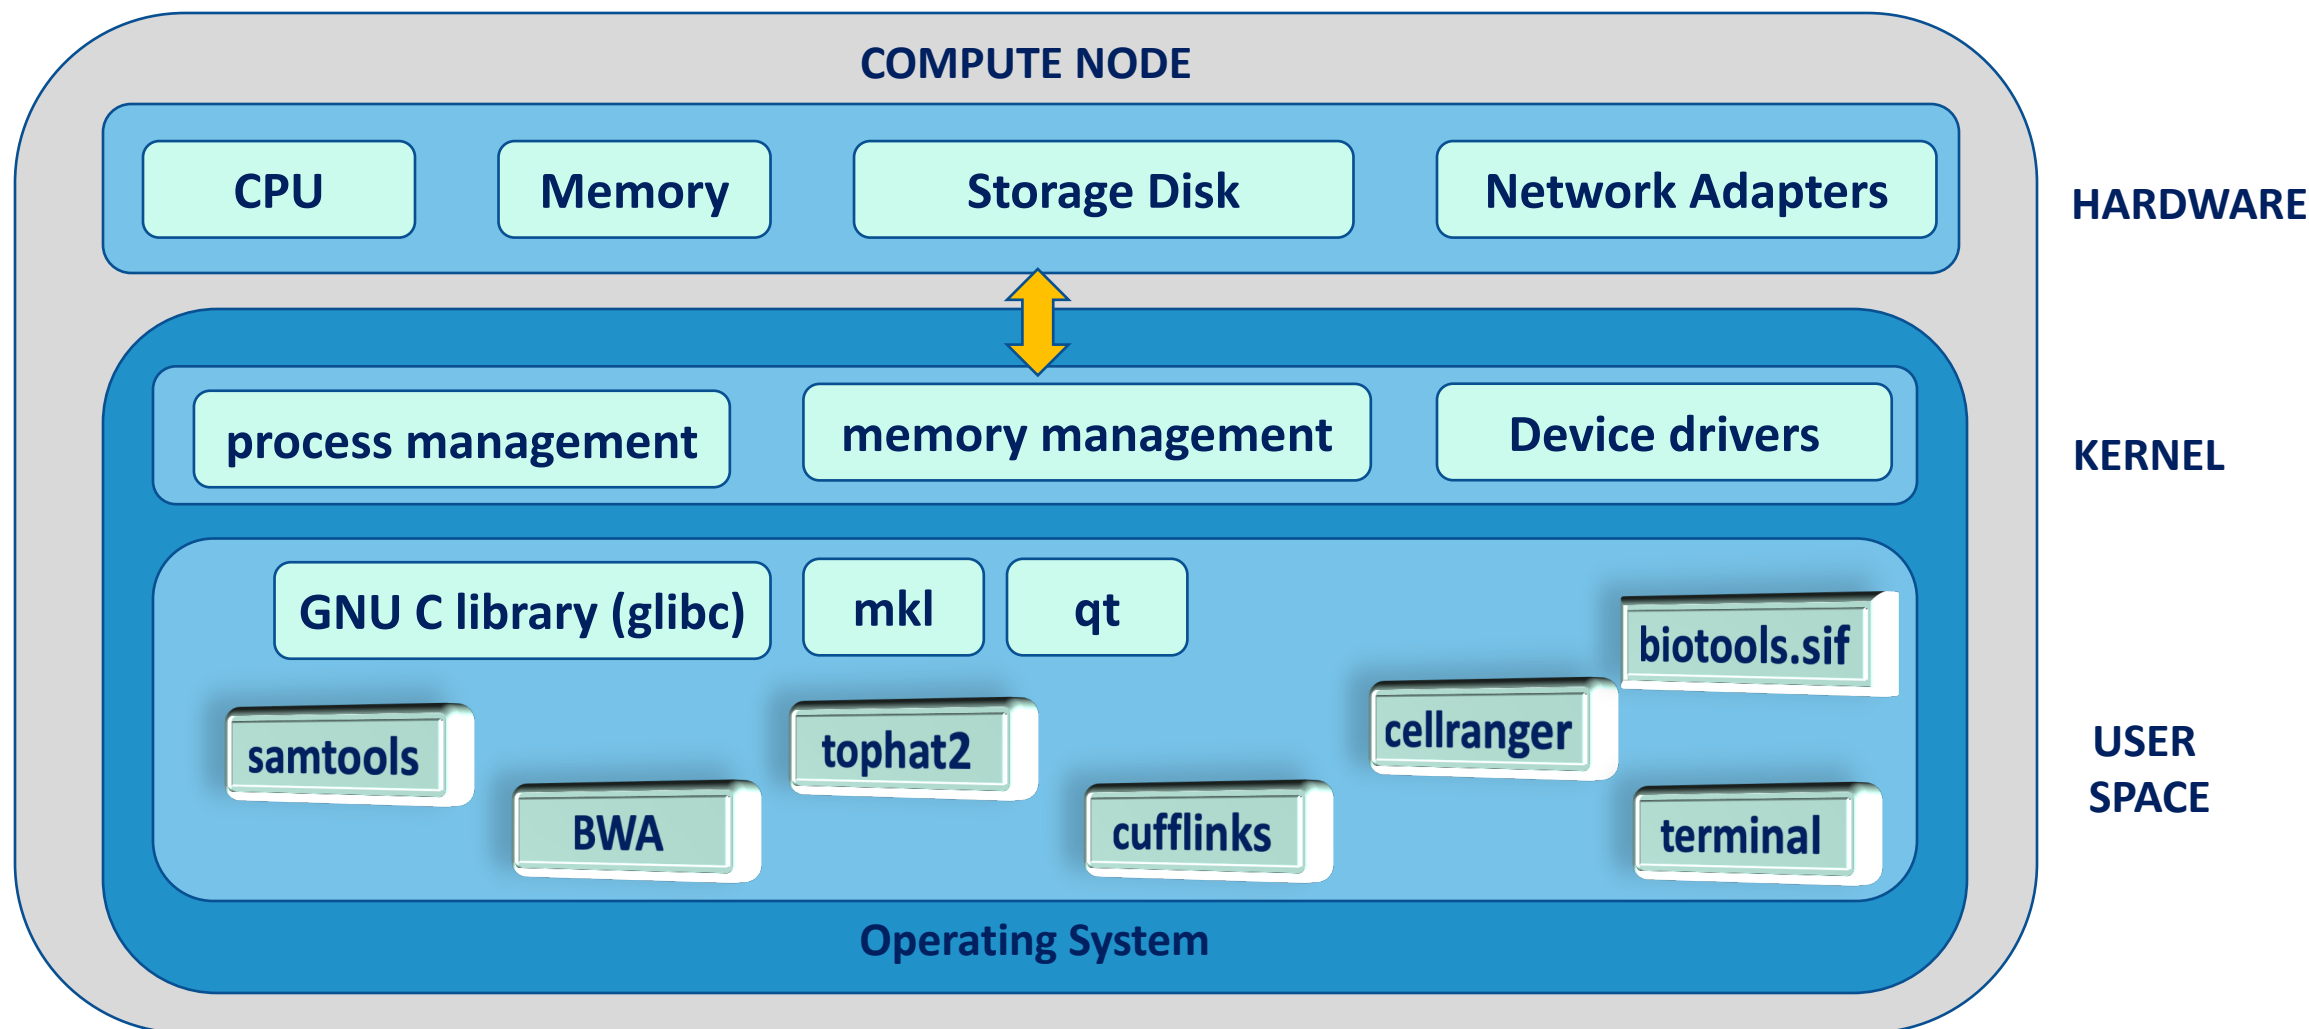
- I. How does High Performance Computing (HPC) environment work?**
 - a. Containerization of Bioinformatics tools
- II. How to get Containers for Bioinformatics tools?**
 - a. Pre-built containers
 - b. Customized containers
- III. How to run containers on a job scheduler**

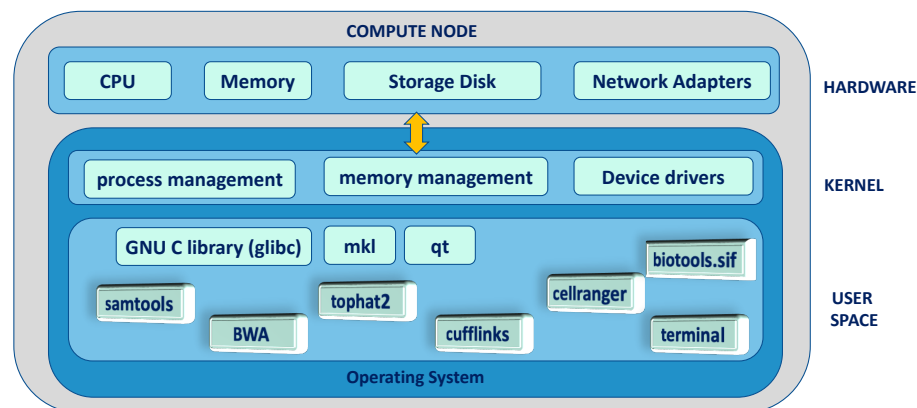
How to get to a HPC environment?

Users on
personal
devices



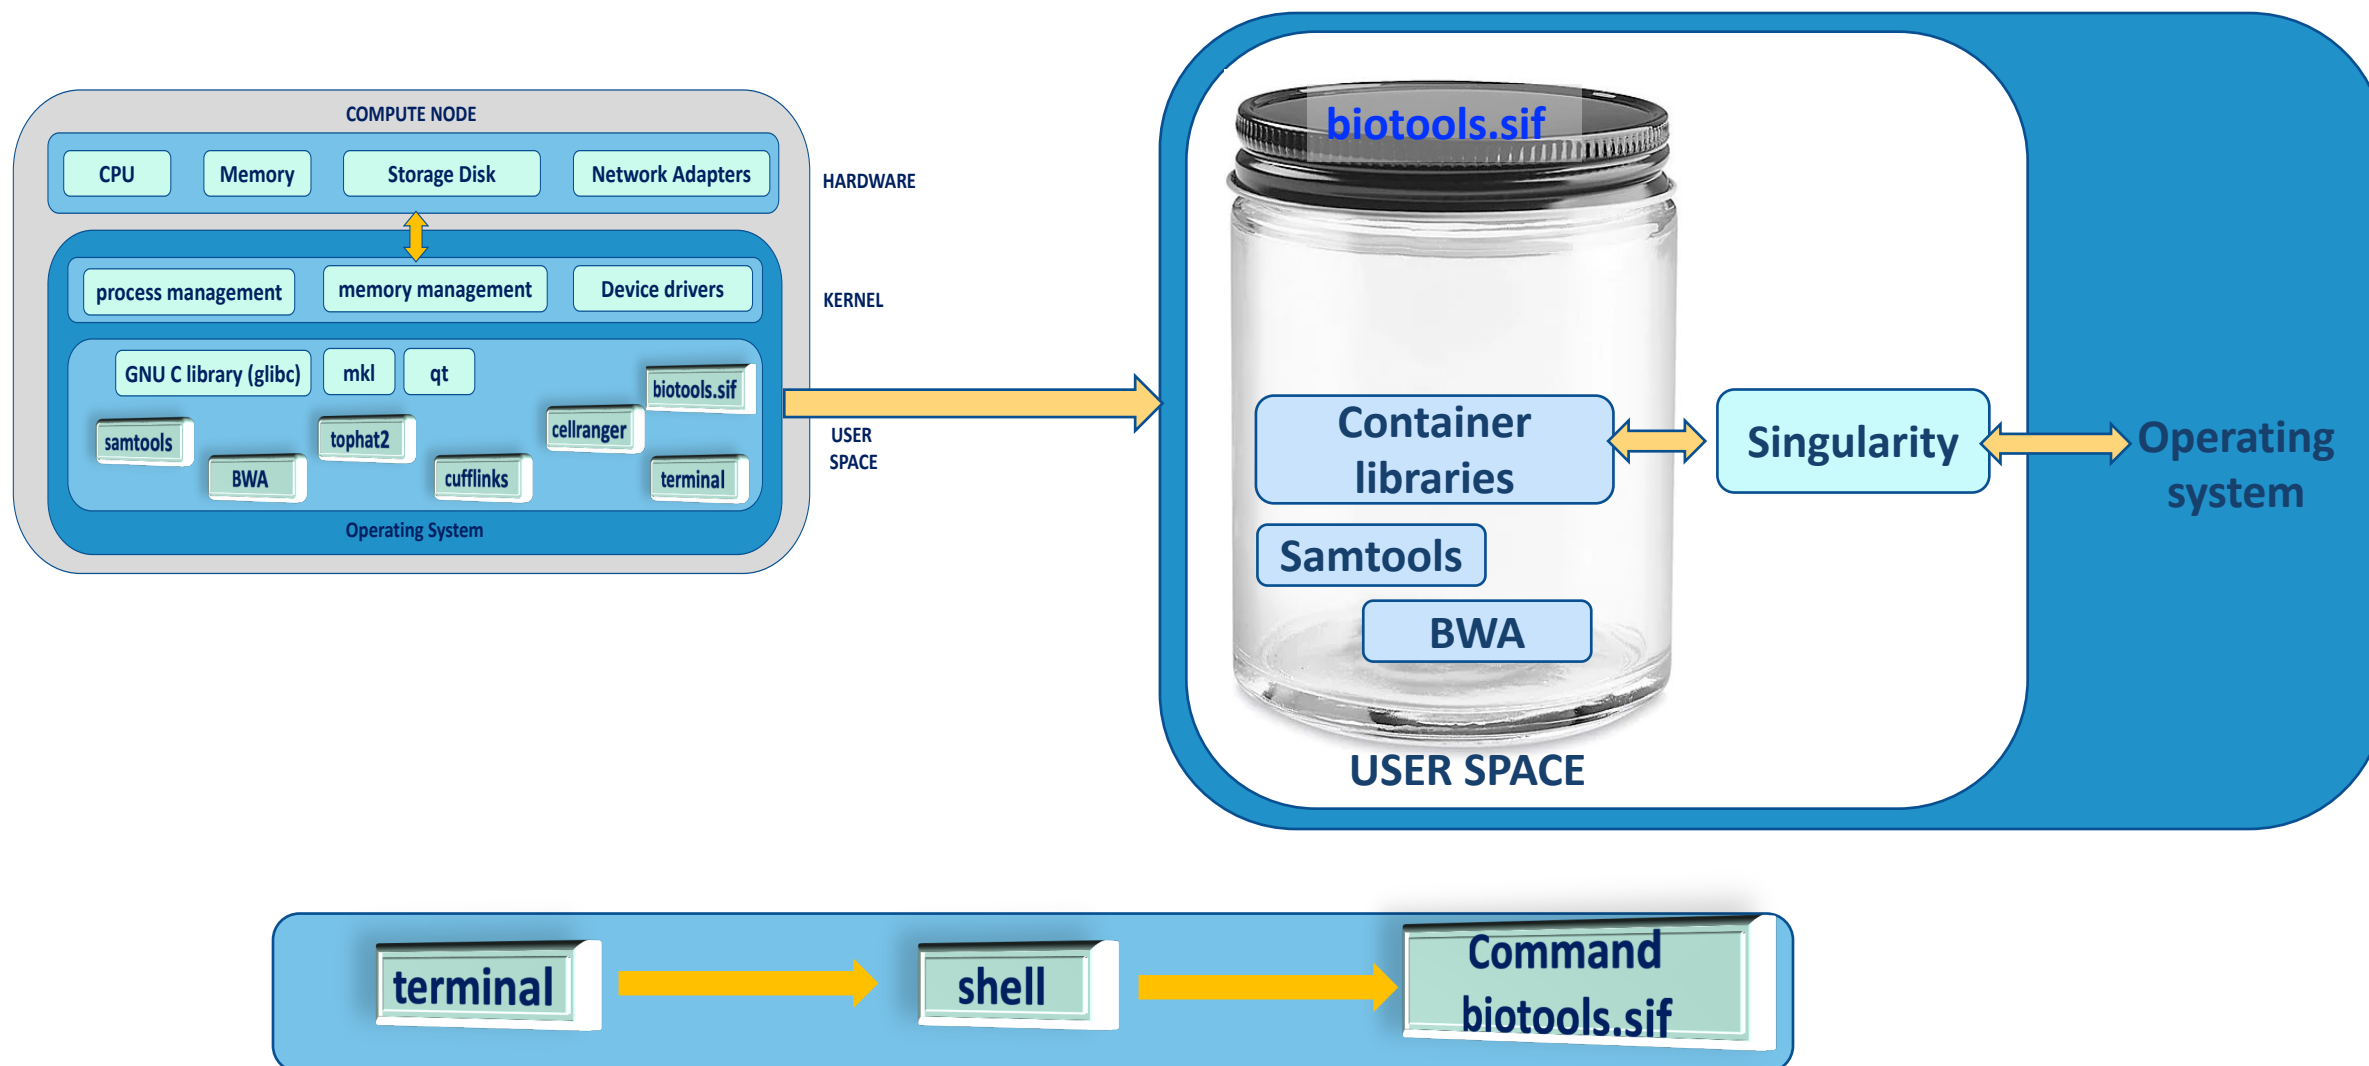
*Example: Compute Canada systems - Cedar or Graham or Beluga
University's local HPC infrastructure*





```
[nandit@cedar1 ~]$ ls /
```

bin	localscratch	project
boot	media	root
cvmfs	misc	run
data	mnt	sbin
dev	nearline	scratch
etc	net	srv
home	opt	sys




Reliable way to capture and distribute software and its compute environments

- Reproducibility
- Portability

Use Case

1. Reproduce results from a published journal article
2. Transfer pipelines from a test environment to a production environment



Reference: Baker, M. *Nature* **533**, 452–454 (2016)

What problems do Containers solve?

Compute Canada systems: the supporting software environment. For example, samtools

CEDAR

```
tannistha.nandi — nandit@cedar1:~ — ssh nandit@cedar.computeCanada.ca — 95x22
[nandit@cedar1 ~]$ module spider samtools

samtools:

Description:
SAM Tools provide various utilities for manipulating alignments in the SAM format,
including sorting, merging, indexing and generating alignments in a per-position
format.

Versions:
samtools/0.1.17
samtools/0.1.18
samtools/0.1.20
samtools/1.3.1
samtools/1.5
samtools/1.8
samtools/1.9
samtools/1.10
samtools/1.11
samtools/1.12
```

BELUGA

```
tannistha.nandi — nandit@beluga3:~ — ssh nandit@beluga.computeCanada.ca — 95x22
[nandit@beluga3 ~]$ module spider samtools

samtools:

Description:
SAM Tools provide various utilities for manipulating alignments in the SAM format,
including sorting, merging, indexing and generating alignments in a per-position
format.

Versions:
samtools/0.1.20
samtools/1.3.1
samtools/1.9
samtools/1.10
samtools/1.11
samtools/1.12

For detailed information about a specific "samtools" package (including how to load the modul
es) use the module's full name.
Note that names that have a trailing (E) are extensions provided by other modules.
```



1. Prebuilt containers

Available public / private repositories or from another researcher

- Singularity Container Library: <https://cloud.sylabs.io/library>
- Docker Hub: <https://hub.docker.com/>
- NGC: <https://ngc.nvidia.com/catalog>
- Biocontainers: <https://biocontainers.pro/>
- Research collaborator

For example:

```
$ singularity pull py39.sif docker://python:3.9
```

Custom name
of the image

Registry
name

2. Customized Container

a. Access to a Local Linux machine with elevated 'sudo' privileges.

- use a device with Linux OS, install Singularity on it.
- use MacBook or Windows laptop/desktop, install a VM to provide Linux OS, then install Singularity in the VM.

b. A recipe file for the singularity container (extension .def)

```
[user@local ~]$ ls  
container.def
```

2. Customized Containers – Example 1: Singularity recipe file

```
[nandit@local ~]$ cat container.def
```

```
Bootstrap: docker  
From: centos:latest
```

HEADER

```
%labels  
    Tannistha Nandi  
%runscript  
    echo "Hello from inside the container..."
```

SECTIONS

- defined by a % character followed by the name of the section.
- All sections are optional.

2. Customized Containers – Example 1: Singularity recipe file

```
[nandit@local ~]$ sudo singularity build container.sif container.def
```

```
[sudo] password for tannistha.nandi:
INFO: Starting build...
Getting image source signatures
Copying blob ald0c7532777 done
Copying config 8c1402b22a done
Writing manifest to image destination
Storing signatures
2021/10/10 08:21:13 info unpack layer: sha256:ald0c75327776413fa0db9ed3adcdbadcdc95a662eb1d360dad82bb913f8a1d1
INFO: Adding labels
INFO: Adding runscript
INFO: Creating SIF file...
INFO: Build complete: container.sif
```

```
[nandit@local ~]$ ls
container.def container.sif
```

```
[nandit@local ~]$ ./container.sif
Hello from inside the container...
```

2. Customized Containers – Example 1: Singularity recipe file

```
[nandit@local ~]$ singularity inspect container.sif
```

```
Tannistha: Nandi
```

```
org.label-schema.build-arch: amd64
```

```
org.label-schema.build-date: Sunday_10_October_2021_8:21:19_MDT
```

```
org.label-schema.schema-version: 1.0
```

```
org.label-schema.usage.singularity.deffile.bootstrap: docker
```

```
org.label-schema.usage.singularity.deffile.from: centos:8
```

```
org.label-schema.usage.singularity.version: 3.8.0-1.el8
```

2. Customized Containers – Example 2: Singularity recipe for a container to run a python script

```
[nandit@local ~]$ ls
python.def myscript.py
[nandit@local ~]$ cat python.def
```

```
Bootstrap: docker
From: python:3.9
%labels
    Tannistha Nandi
%files
    /path/to/myscript.py /opt/
%post
    apt-get -y update
    apt-get -y install pip wget git
%runscript
    python myscript.py
```

2. Customized Containers – Example 2: Singularity recipe for a container to run a python script

a. Explore the environment within Singularity shell

```
[nandit@local ~]$ singularity shell python.sif
Singularity>
Singularity> ls /opt
myscript.py
Singularity> exit
```

b. Execute the python program

```
[nandit@local ~]$ singularity exec python.sif python myscript.py
This is my first python script. I am happy!
[nandit@local ~]$
```

c. Run the default “runscript” in the container

```
[nandit@local ~]$ singularity run python.sif
This is my first python script. I am happy!
```

2. Customized Containers – Example 3: Singularity recipe for ‘samtools’ container

%runscript: Define commands that will be executed by singularity run.

%post: Execute commands after the base OS has been installed

```

Bootstrap: docker
From: centos:8

%labels
    Tannistha Nandi
%runscript
    exec "@@"
%setup
    mkdir -p ${SINGULARITY_ROOTFS}/opt
    mkdir -p ${SINGULARITY_ROOTFS}/scratch
    mkdir -p ${SINGULARITY_ROOTFS}/shared
%post
    dnf -y makecache
    dnf -y group install "Development Tools"
    dnf -y install --allowerasing hostname which dnf-utils git zlib zlib-devel \
        bzip2 bzip2-devel xz xz-devel libcurl libcurl-devel ncurses ncurses-devel \
        unzip wget gnuplot rsync java-1.8.0-openjdk java-1.8.0-openjdk-devel \
        openssl-devel libffi-devel
%environment
    export PATH=$PATH:/usr/bin:/usr/local/bin/
    LANG="en_US.UTF-8"
    
```


2. Customized Containers – Example 3: Singularity recipe for ‘samtools’ container

%app sections - build a single container with two or three different programs*

```
%apprun samtools
```

```
exec samtools
```


```
%appinstall samtools
```

```
git clone https://github.com/samtools/htslib
cd htslib
autoheader
autoconf
git submodule update --init --recursive
./configure
make -j 8
make install
cd ..
git clone https://github.com/samtools/samtools
cd samtools
autoheader
autoconf -Wno-syntax
./configure
make -j 8
make install
cd ..
```

%apprun: runs script for the app

%appinstall: commands to install the app (similar to *%post* but only for one app here).

Run containers on a job scheduler

`[nandit@local ~]$`  `[nandit@cedar ~]$`
`biotools.sif` **transfer the container image to the HPC system** `biotools.sif`

Container image is a file with .sif or .sing extension that contains everything needed to run applications.

Run the container **without elevated privilege** on the host HPC system like cedar.

Use containers

1. native command

```
$ ./biotools.sif
```

2. subcommand 'shell'

```
$ singularity shell biotools.sif
```

3. subcommand 'exec'

```
$ singularity exec biotools.sif samtools
```

4. subcommand 'run'

```
$ singularity run biotools.sif
```

Run containers on a job scheduler

1. Get an interactive session

```
[nandit@cedar1~]$ salloc --account=def-<user> --node=1 --mem=6G \  
--time=02:00:00 --cpus-per-task=2
```

```
[nandit@cdr767 ~]$
```

```
[nandit@cdr767~]$ module load singularity
```

```
[nandit@cdr767~]$ singularity shell biotools.sif
```

```
Singularity> samtools --help
```

```
Program: samtools (Tools for alignments in the SAM format)
```

```
Version: 1.13-14-g65a97fe (using htslib 1.13-19-g31bf087)
```

```
Usage: samtools <command> [options]
```

Run containers on a job scheduler

2. Batch jobs

A job script (saved as `jobscript.sh`) to run the `samtools` program using Singularity

```
[nandit@cedar1~]$ cat jobscript.sh
```

```
#!/bin/bash
```

```
#SBATCH --account=def-<user>
```

```
#SBATCH --node=1
```

```
#SBATCH --time=02:00:00
```

```
#SBATCH --mem=6G
```

```
#SBATCH --cpus-per-task=2
```

```
module load singularity
```

```
singularity exec -B /path/to/input:/data biotools.sif samtools index -bc -@ 2  
/data/input.bam input.index
```

Submit the job script

```
[nandit@cedar1~]$ sbatch jobscript.sh
```



**Happy to take your
questions!!**

Thank you for your attention.....