

# TENSORBOARD & OBJECT DETECTION

12 OCTOBER 2022  
JILLIAN ANDERSON



# OVERVIEW

- The Individual Pieces
  - Object detection
  - TensorFlow
  - TensorBoard
- TensorBoard + Object Detection
- TensorBoard on HPC
- Understanding the TensorBoard interface

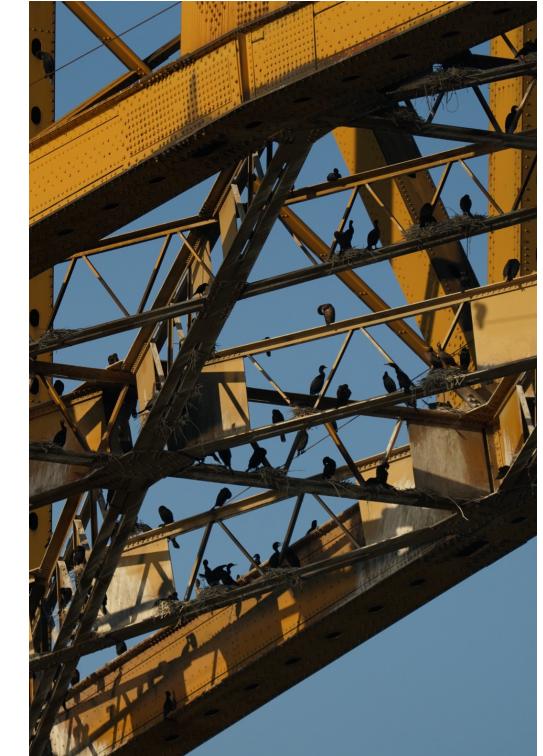
# OBJECT DETECTION

# OBJECT DETECTION

- Computer vision task where a trained model identify objects within an image or video
- Applications include self-driving vehicles, remote sensing, wildlife monitoring



# WILDLIFE MONITORING EXAMPLE



# TENSORFLOW

# TENSORFLOW

- A free & open source library for machine learning, focused on deep neural networks
- The Alliance [TensorFlow documentation](#)

\* TF is available in other languages (e.g. JavaScript, C++, Java) but is most commonly used with Python

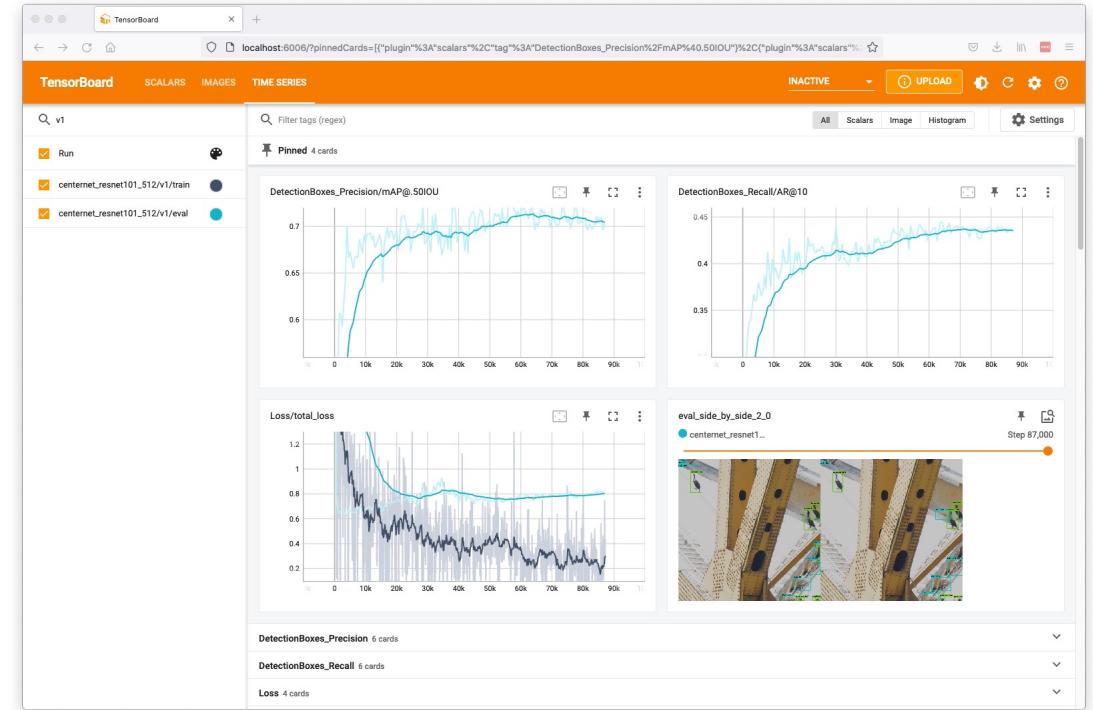
# OBJECT DETECTION & TENSORFLOW

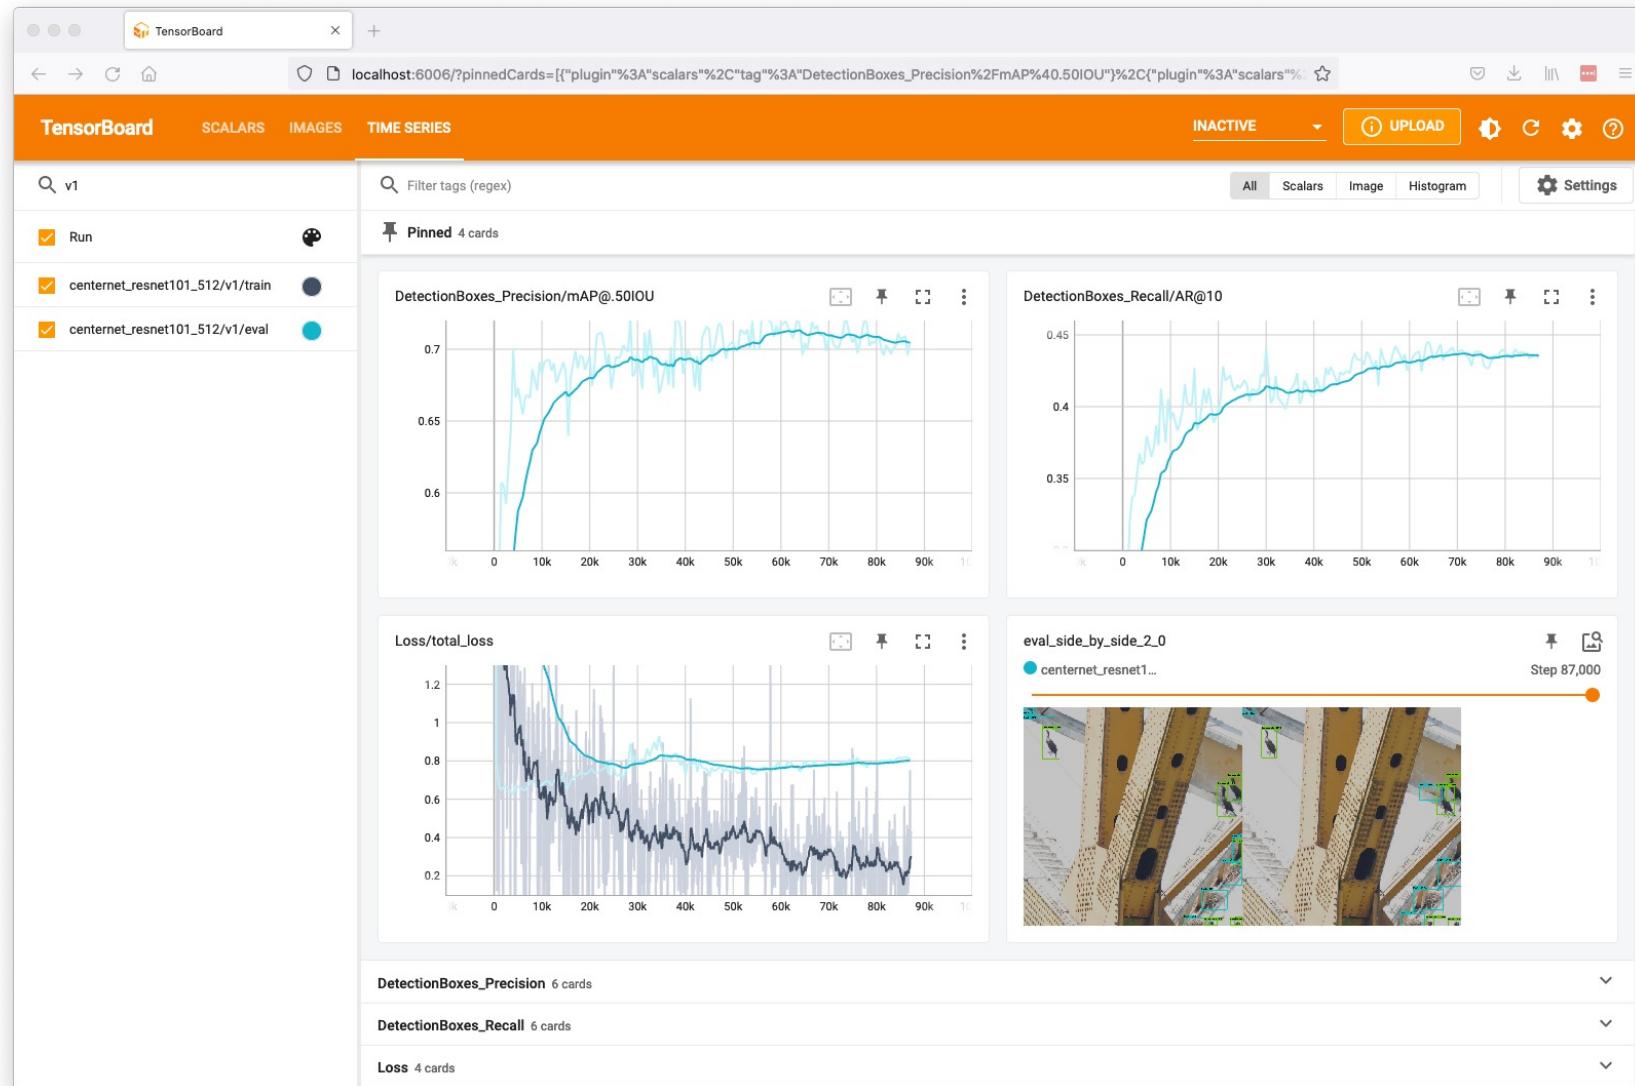
- The TensorFlow Object Detection API is one of the research models/implementations provided in the [TF Model Garden](#)
- Learn more in TF's [object detection tutorial](#) & the Object Detection API's [Training & Evaluation with TensorFlow 2](#) guide

# TENSORBOARD

# WHAT IS TENSORBOARD?

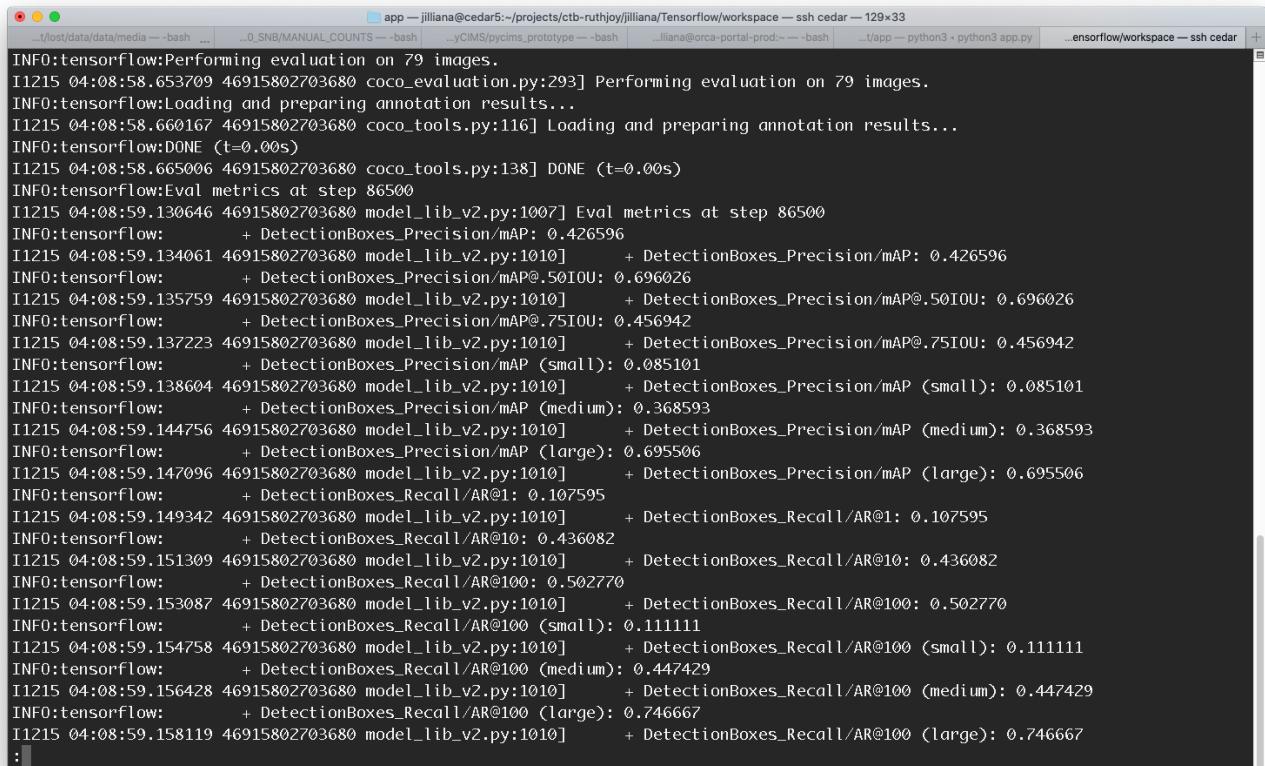
- TF's visual toolkit for machine learning experimentation
- Visually track model information (e.g. model structure, mAP, etc)





# WHY USE TENSORBOARD?

- Machine Learning usually requires experimentation
- We need a way to compare models (aka runs)
- Ideally we want to avoid manually inspecting logs



The screenshot shows a terminal window with multiple tabs. The active tab displays TensorFlow evaluation logs. The logs show the process of performing evaluation on 79 images, loading annotation results, and calculating various metrics like mAP and recall at different IoU thresholds (0.50, 0.75, and large). The logs are timestamped from 04:08:58 to 04:08:59.

```
INFO:tensorflow:Performing evaluation on 79 images.
I1215 04:08:58.653709 46915802703680 coco_evaluation.py:293] Performing evaluation on 79 images.
INFO:tensorflow:Loading and preparing annotation results...
I1215 04:08:58.660167 46915802703680 coco_tools.py:116] Loading and preparing annotation results...
INFO:tensorflow:DONE (t=0.00s)
I1215 04:08:58.665006 46915802703680 coco_tools.py:138] DONE (t=0.00s)
INFO:tensorflow:Eval metrics at step 86500
I1215 04:08:59.130646 46915802703680 model_lib_v2.py:1007] Eval metrics at step 86500
INFO:tensorflow:    + DetectionBoxes_Precision/mAP: 0.426596
I1215 04:08:59.134061 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Precision/mAP: 0.426596
INFO:tensorflow:    + DetectionBoxes_Precision/mAP@.50IoU: 0.696026
I1215 04:08:59.135759 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Precision/mAP@.50IoU: 0.696026
INFO:tensorflow:    + DetectionBoxes_Precision/mAP@.75IoU: 0.456942
I1215 04:08:59.137223 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Precision/mAP@.75IoU: 0.456942
INFO:tensorflow:    + DetectionBoxes_Precision/mAP (small): 0.085101
I1215 04:08:59.138604 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Precision/mAP (small): 0.085101
INFO:tensorflow:    + DetectionBoxes_Precision/mAP (medium): 0.368593
I1215 04:08:59.144756 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Precision/mAP (medium): 0.368593
INFO:tensorflow:    + DetectionBoxes_Precision/mAP (large): 0.695506
I1215 04:08:59.147096 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Precision/mAP (large): 0.695506
INFO:tensorflow:    + DetectionBoxes_Recall/AR@1: 0.107595
I1215 04:08:59.149342 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Recall/AR@1: 0.107595
INFO:tensorflow:    + DetectionBoxes_Recall/AR@10: 0.436082
I1215 04:08:59.151309 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Recall/AR@10: 0.436082
INFO:tensorflow:    + DetectionBoxes_Recall/AR@100: 0.502770
I1215 04:08:59.153087 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Recall/AR@100: 0.502770
INFO:tensorflow:    + DetectionBoxes_Recall/AR@100 (small): 0.111111
I1215 04:08:59.154758 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Recall/AR@100 (small): 0.111111
INFO:tensorflow:    + DetectionBoxes_Recall/AR@100 (medium): 0.447429
I1215 04:08:59.156428 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Recall/AR@100 (medium): 0.447429
INFO:tensorflow:    + DetectionBoxes_Recall/AR@100 (large): 0.746667
I1215 04:08:59.158119 46915802703680 model_lib_v2.py:1010]    + DetectionBoxes_Recall/AR@100 (large): 0.746667
:
```

# HOW DOES TENSORBOARD WORK?

- TensorFlow writes event files during model training
- TensorBoard parses these files & generates visualizations
- You can control what events are logged using FileWriter instance(s)

# MORE ON TENSORBOARD

→ [TensorBoard guide](#)

The screenshot shows a web browser displaying the TensorFlow website at [https://www.tensorflow.org/tensorboard/get\\_started](https://www.tensorflow.org/tensorboard/get_started). The page title is "Get started with TensorBoard". The navigation bar includes links for "Install", "Learn", "Resources", "More", "Search", "English", "GitHub", and "Sign in". The main content area is titled "TensorBoard" and has tabs for "Overview" and "Guide". The "Guide" tab is selected. A sidebar on the left lists various "Get started" tools: Scalars and metrics, Image data, Model graphs, Text data, Hyperparameter tuning, Embedding projector, What-If tool, Fairness Indicators, Profiling tool, TensorBoard in notebooks, TensorBoard Data as DataFrames, and Debugger V2. The main content area contains a heading "Get started with TensorBoard", a "On this page" section with links to "Using TensorBoard with Keras Model.fit()", "Using TensorBoard with other methods", and "TensorBoard.dev: Host and share your ML experiment results". It also features three buttons: "Run in Google Colab", "View source on GitHub", and "Download notebook". A descriptive paragraph at the bottom explains that TensorBoard is a tool for tracking experiment metrics like loss and accuracy, visualizing the model graph, projecting embeddings to a lower dimensional space, and more.

# TENSORBOARD + OBJECT DETECTION

# SETTING UP EVENT FILES

- Setup pipeline.config to generate event files for **scalars\*** and **images\***
- eval\_config – how should evaluation be done?
- eval\_input\_reader – what data should be used for evaluation?

```
eval_config: {  
    metrics_set: "coco_detection_metrics"  
    use_moving_averages: false  
    batch_size: 1  
    num_visualizations: 50  
    min_score_threshold: 0.1  
}  
  
eval_input_reader: {  
    label_map_path: "models/snb3/label_map.pbtxt"  
    shuffle: false  
    num_epochs: 1  
    tf_record_input_reader {  
        input_path: "models/snb3/validation.tfrecord-*"  
    }  
}
```

\* Other info (e.g. graph structure, Tensor histograms) can be tracked too, but we won't cover this today.

# SCALARS

→ By default evaluation uses  
COCO detection metrics

## 2. Metrics

The following 12 metrics are used for characterizing the performance of an object detector on COCO:

<b>Average Precision (AP):</b>	
AP	% AP at IoU=.50:.05:.95 ( <b>primary challenge metric</b> )
AP <sup>IoU=.50</sup>	% AP at IoU=.50 (PASCAL VOC metric)
AP <sup>IoU=.75</sup>	% AP at IoU=.75 (strict metric)
<b>AP Across Scales:</b>	
AP <sup>small</sup>	% AP for small objects: area < 32 <sup>2</sup>
AP <sup>medium</sup>	% AP for medium objects: 32 <sup>2</sup> < area < 96 <sup>2</sup>
AP <sup>large</sup>	% AP for large objects: area > 96 <sup>2</sup>
<b>Average Recall (AR):</b>	
AR <sup>max=1</sup>	% AR given 1 detection per image
AR <sup>max=10</sup>	% AR given 10 detections per image
AR <sup>max=100</sup>	% AR given 100 detections per image
<b>AR Across Scales:</b>	
AR <sup>small</sup>	% AR for small objects: area < 32 <sup>2</sup>
AR <sup>medium</sup>	% AR for medium objects: 32 <sup>2</sup> < area < 96 <sup>2</sup>
AR <sup>large</sup>	% AR for large objects: area > 96 <sup>2</sup>

1. Unless otherwise specified, AP and AR are averaged over multiple Intersection over Union (IoU) values. Specifically we use 10 IoU thresholds of .50:.05:.95. This is a break from tradition, where AP is computed at a single IoU of .50 (which corresponds to our metric AP<sup>IoU=.50</sup>). Averaging over IoUs rewards detectors with better localization.
2. AP is averaged over all categories. Traditionally, this is called "mean average precision" (mAP). We make no distinction between AP and mAP (and likewise AR and mAR) and assume the difference is clear from context.
3. AP (averaged across all 10 IoU thresholds and all 80 categories) will determine the challenge winner. This should be considered the single most important metric when considering performance on COCO.
4. In COCO, there are more small objects than large objects. Specifically: approximately 41% of objects are small (area < 32<sup>2</sup>), 34% are medium (32<sup>2</sup> < area < 96<sup>2</sup>), and 24% are large (area > 96<sup>2</sup>). Area is measured as the number of pixels in the segmentation mask.
5. AR is the maximum recall given a fixed number of detections per image, averaged over categories and IoUs. AR is related to the metric of the same name used in [proposal evaluation](#) but is computed on a per-category basis.
6. All metrics are computed allowing for at most 100 top-scoring detections per image (across all categories).
7. The evaluation metrics for detection with bounding boxes and segmentation masks are identical in all respects except for the IoU computation (which is performed over boxes or masks, respectively).

# SCALARS

- Specified using `metrics_set`
- See the [docs](#) for the other available metrics

```
eval_config: {
    metrics_set: "coco_detection_metrics"
    use_moving_averages: false
    batch_size: 1
    num_visualizations: 50
    min_score_threshold: 0.1
}

eval_input_reader: {
    label_map_path: "models/snb3/label_map.pbtxt"
    shuffle: false
    num_epochs: 1
    tf_record_input_reader {
        input_path: "models/snb3/validation.tfrecord-*"
    }
}
```

# IMAGES

- Visualize model predictions on evaluation data
- Compare ground truth with predictions

```
eval_config: {  
    metrics_set: "coco_detection_metrics"  
    use_moving_averages: false  
    batch_size: 1  
    num_visualizations: 50  
    min_score_threshold: 0.1  
}  
  
eval_input_reader: {  
    label_map_path: "models/snb3/label_map.pbtxt"  
    shuffle: false  
    num_epochs: 1  
    tf_record_input_reader {  
        input_path: "models/snb3/validation.tfrecord-*"  
    }  
}
```

# IMAGES

- Specified with `num_visualizations`
- `min_score_threshold` controls which detections are visualized
- To explore other options, checkout the [eval.proto](#) file

```
eval_config: {  
    metrics_set: "coco_detection_metrics"  
    use_moving_averages: false  
    batch_size: 1  
    num_visualizations: 50  
    min_score_threshold: 0.1  
}  
  
eval_input_reader: {  
    label_map_path: "models/snb3/label_map.pbtxt"  
    shuffle: false  
    num_epochs: 1  
    tf_record_input_reader {  
        input_path: "models/snb3/validation.tfrecord-*"  
    }  
}
```

# TENSORBOARD ON HPC

# TENSORBOARD ON HPC

- During model training
- After model training is complete

# DURING MODEL TRAINING

- Adjust your job script
- Submit your job as normal
- Connect to the node running the job
- Visit TensorBoard in your browser

# DURING MODEL TRAINING

**ADJUST → SUBMIT → CONNECT → VISIT**

→ Activate evaluation in your job script

```
python model_main_tf2.py \
  --pipeline_config_path=$MODELDIR/pipeline.config \
  --model_dir=$MODELDIR \
  --checkpoint_dir=$MODELDIR \
  --num_workers=1 \
  --sample_1_of_n_eval_examples=1 &
```

→ Activate TensorBoard in your job script

```
tensorboard --logdir=$MODELDIR --host 0.0.0.0 --load_fast false &
```

# DURING MODEL TRAINING

**ADJUST → SUBMIT → CONNECT → VISIT**

- Submit your job to the cluster as normal  
`sbatch train_model.sh`
- Wait for your job to begin

# DURING MODEL TRAINING

ADJUST → SUBMIT → CONNECT → VISIT

→ Find what node your model is running on  
squeue -u user

JOBID	USER	ACCOUNT	NAME	ST	TIME_LEFT	NODES	CPUS	TRES_PER_N	MIN_MEM	NODELIST	(REASON)
47623590	jilliana	def-jilliana	train.sh	R	9:43	1	1	gres:gpu:1	32G	cdr250	(None)

↑  
Job ID

↑  
HPC  
node

# DURING MODEL TRAINING

ADJUST → SUBMIT → **CONNECT** → VISIT

- Find the node your model is running on
  - Find the port TensorBoard is running on  
grep "TensorBoard" slurm-JOBID.out
- TensorBoard** 2.6.0 at <http://0.0.0.0:6006/> (Press CTRL+C to quit)
-   
↑  
Port  
on  
HPC  
node

# DURING MODEL TRAINING

**ADJUST → SUBMIT → CONNECT → VISIT**

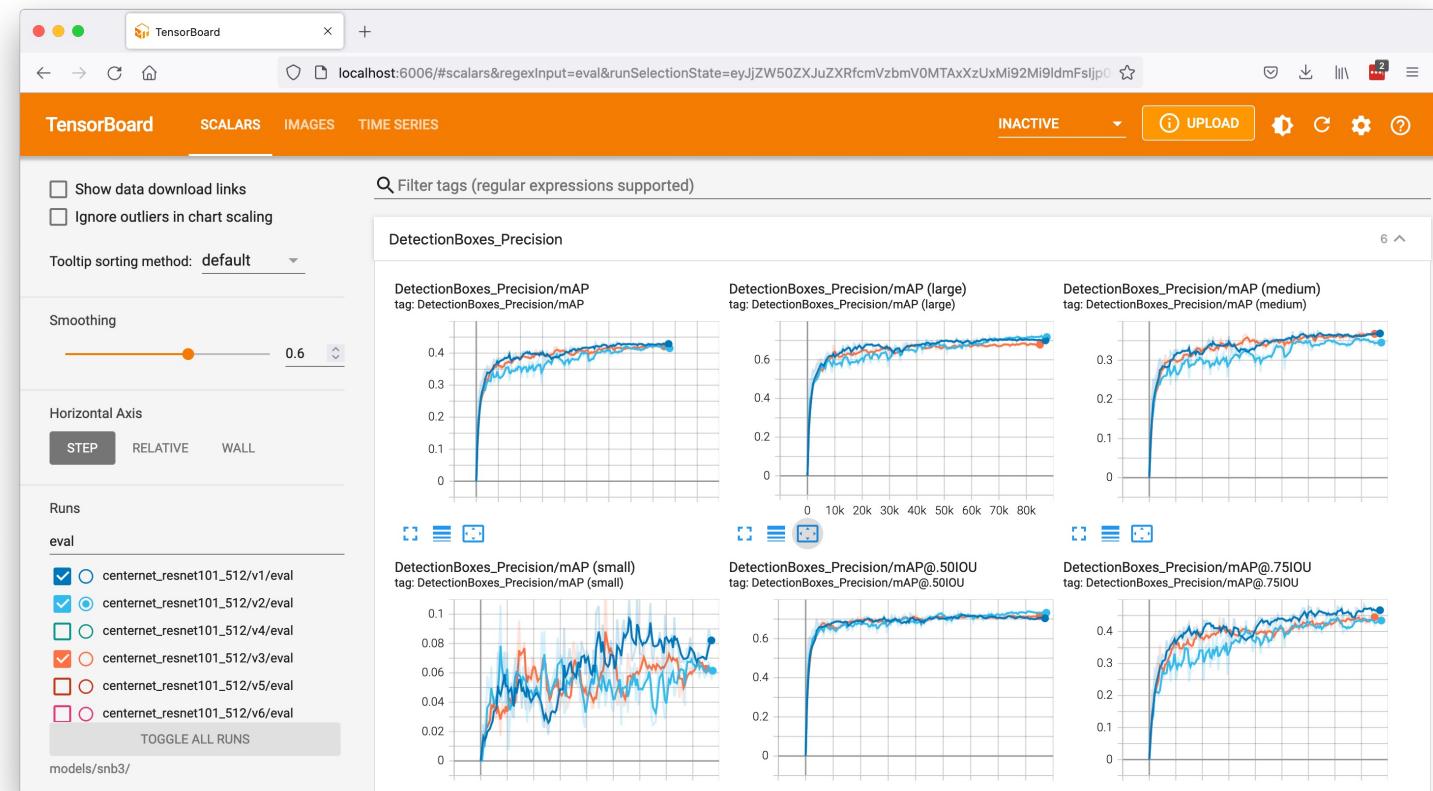
- Find what node your model is running on
- Find what port TensorBoard is running on
- Connect to the node where training is happening  
`ssh -N -f -L localhost:6006:cdr917:6006 user@cedar.computecanada.ca`

HPC node  
Port on HPC node  
Your user  
Your HPC cluster

# DURING MODEL TRAINING

ADJUST → SUBMIT → CONNECT → **VISIT**

→ Visit [localhost:6006](http://localhost:6006)  
on your favourite  
browser



# AFTER MODEL TRAINING\*

- Request an interactive job
- Start TensorBoard
- Connect to the node running the job
- Visit TensorBoard in your browser

\* Requires event files to have been generated during training

# AFTER MODEL TRAINING

**REQUEST** → START TENSORBOARD → CONNECT → VISIT

→ Request an interactive job

```
salloc --mem=16GB --time=1:00:00
```

→ Wait for the interactive session to begin

```
[jilliana@cedar1 Tensorflow]$ salloc --account=def-jilliana --time=1:00:00
salloc: Pending job allocation 47519849
salloc: job 47519849 queued and waiting for resources
salloc: job 47519849 has been allocated resources
salloc: Granted job allocation 47519849
[jilliana@cdr544 Tensorflow]$
```



HPC  
node

# AFTER MODEL TRAINING

REQUEST → START TENSORBOARD → CONNECT → VISIT

## Run TensorBoard

```
tensorboard --logdir=models/dir/ --host 0.0.0.0 --load_fast false
```

## Wait for confirmation that TensorBoard is running

```
[jilliana@cdr783 workspace]$ tensorboard --logdir=models/snb3/ --host 0.0.0.0 --load_fast false
2022-10-10 15:49:15.936765: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'libcud
a.so.1'; dlerror: libcuda.so.1: cannot open shared object file: No such file or directory
2022-10-10 15:49:15.937499: W tensorflow/stream_executor/cuda/cuda_driver.cc:269] failed call to cuInit: UNKNOWN ERROR (303)
2022-10-10 15:49:15.937903: I tensorflow/stream_executor/cuda/cuda_diagnostics.cc:156] kernel driver does not appear to be running
on this host (cdr783.int.cedar.computecanada.ca): /proc/driver/nvidia/version does not exist
TensorBoard 2.9.1 at http://0.0.0.0:6006/ (Press CTRL+C to quit)
```



Port on  
HPC node

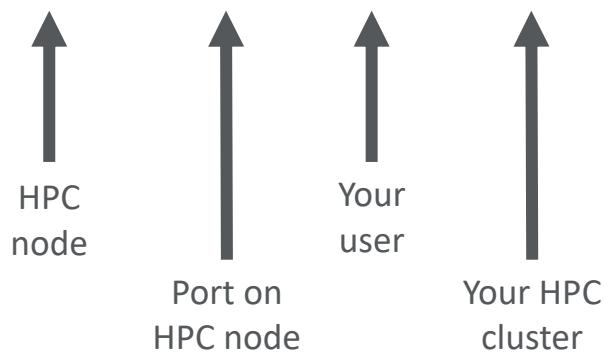
# AFTER MODEL TRAINING

REQUEST → START TENSORBOARD → CONNECT → VISIT

→ Find what node & port your job is running on

→ Connect to the node where training is happening

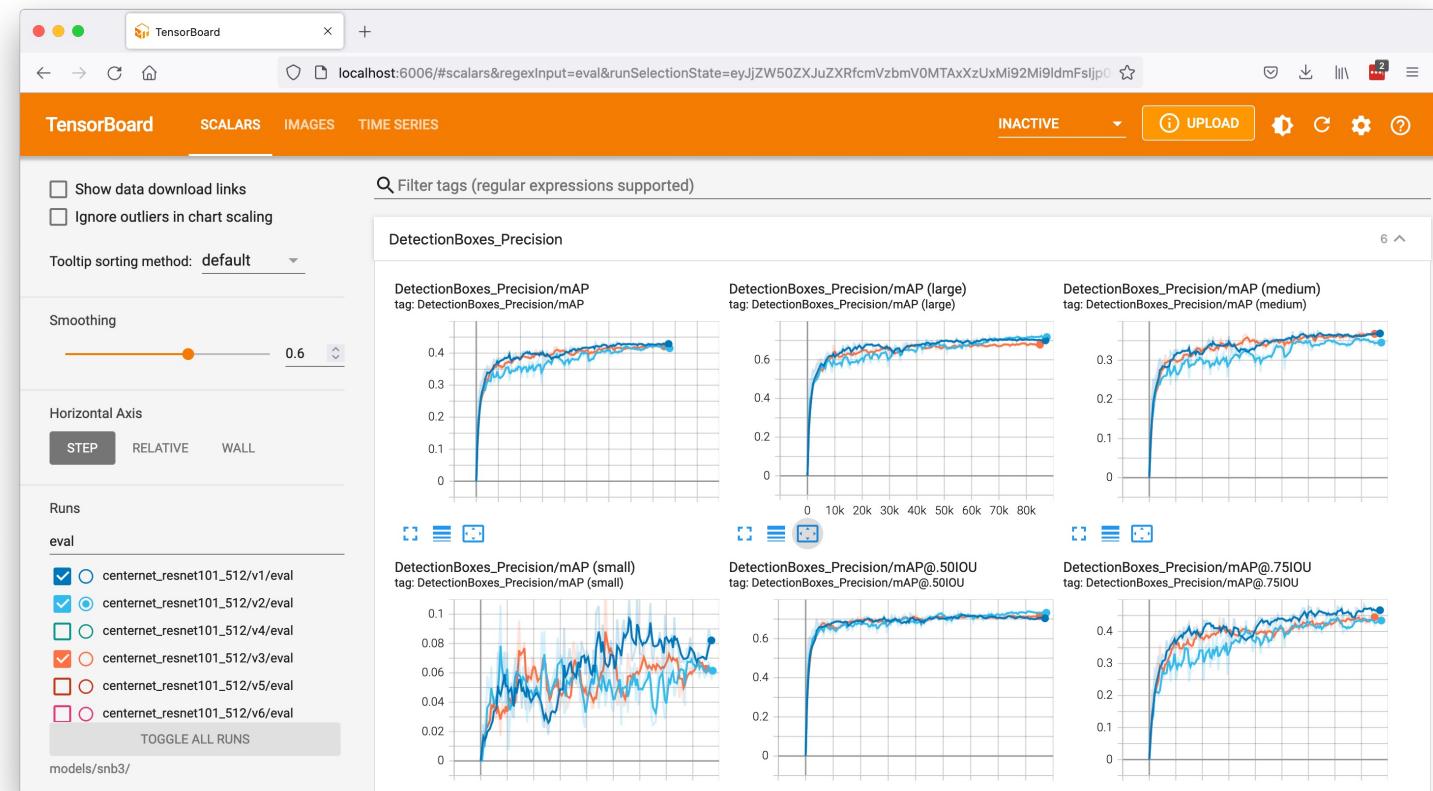
```
ssh -N -f -L localhost:6006:cdr917:6006 user@cedar.computecanada.ca
```



# AFTER MODEL TRAINING

REQUEST → START TENSORBOARD → CONNECT → VISIT

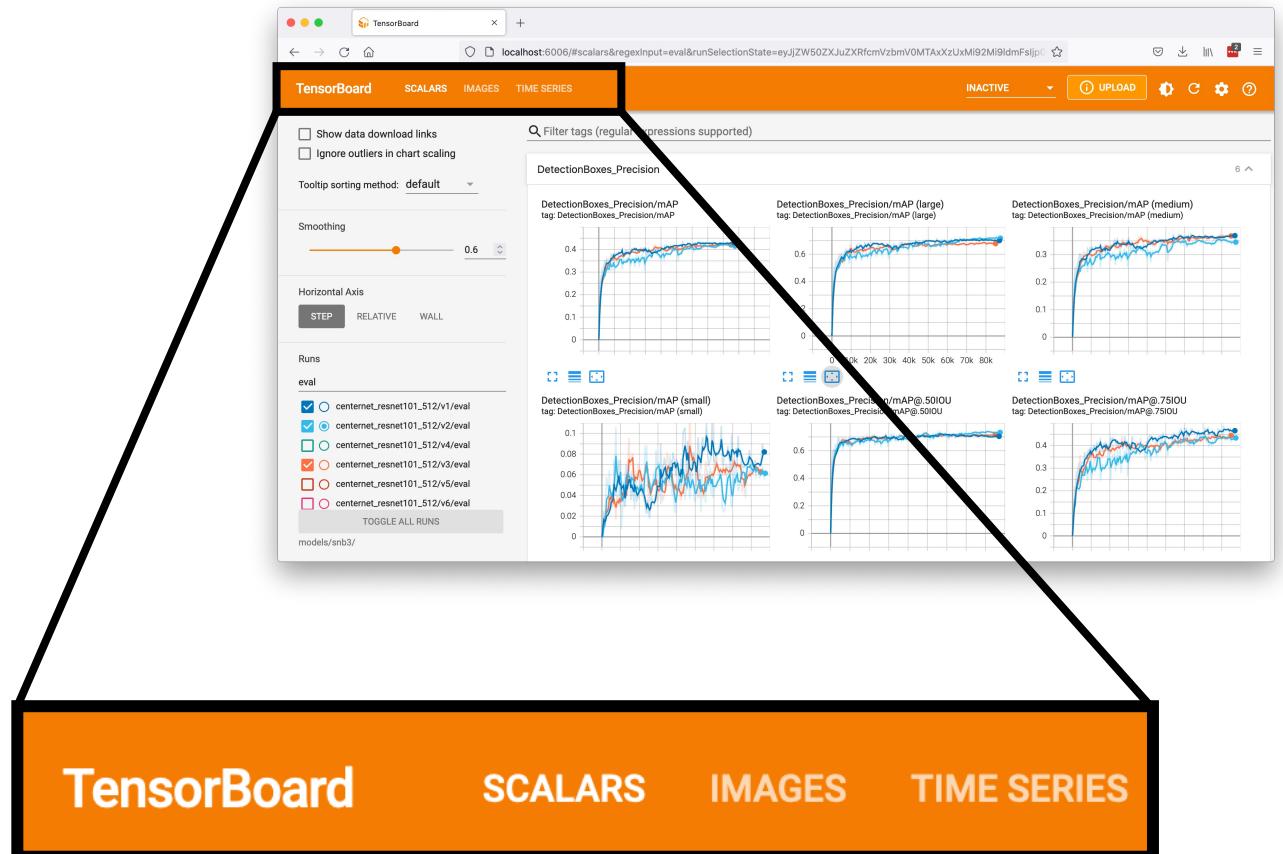
→ Visit [localhost:6006](http://localhost:6006)  
on your favourite  
browser



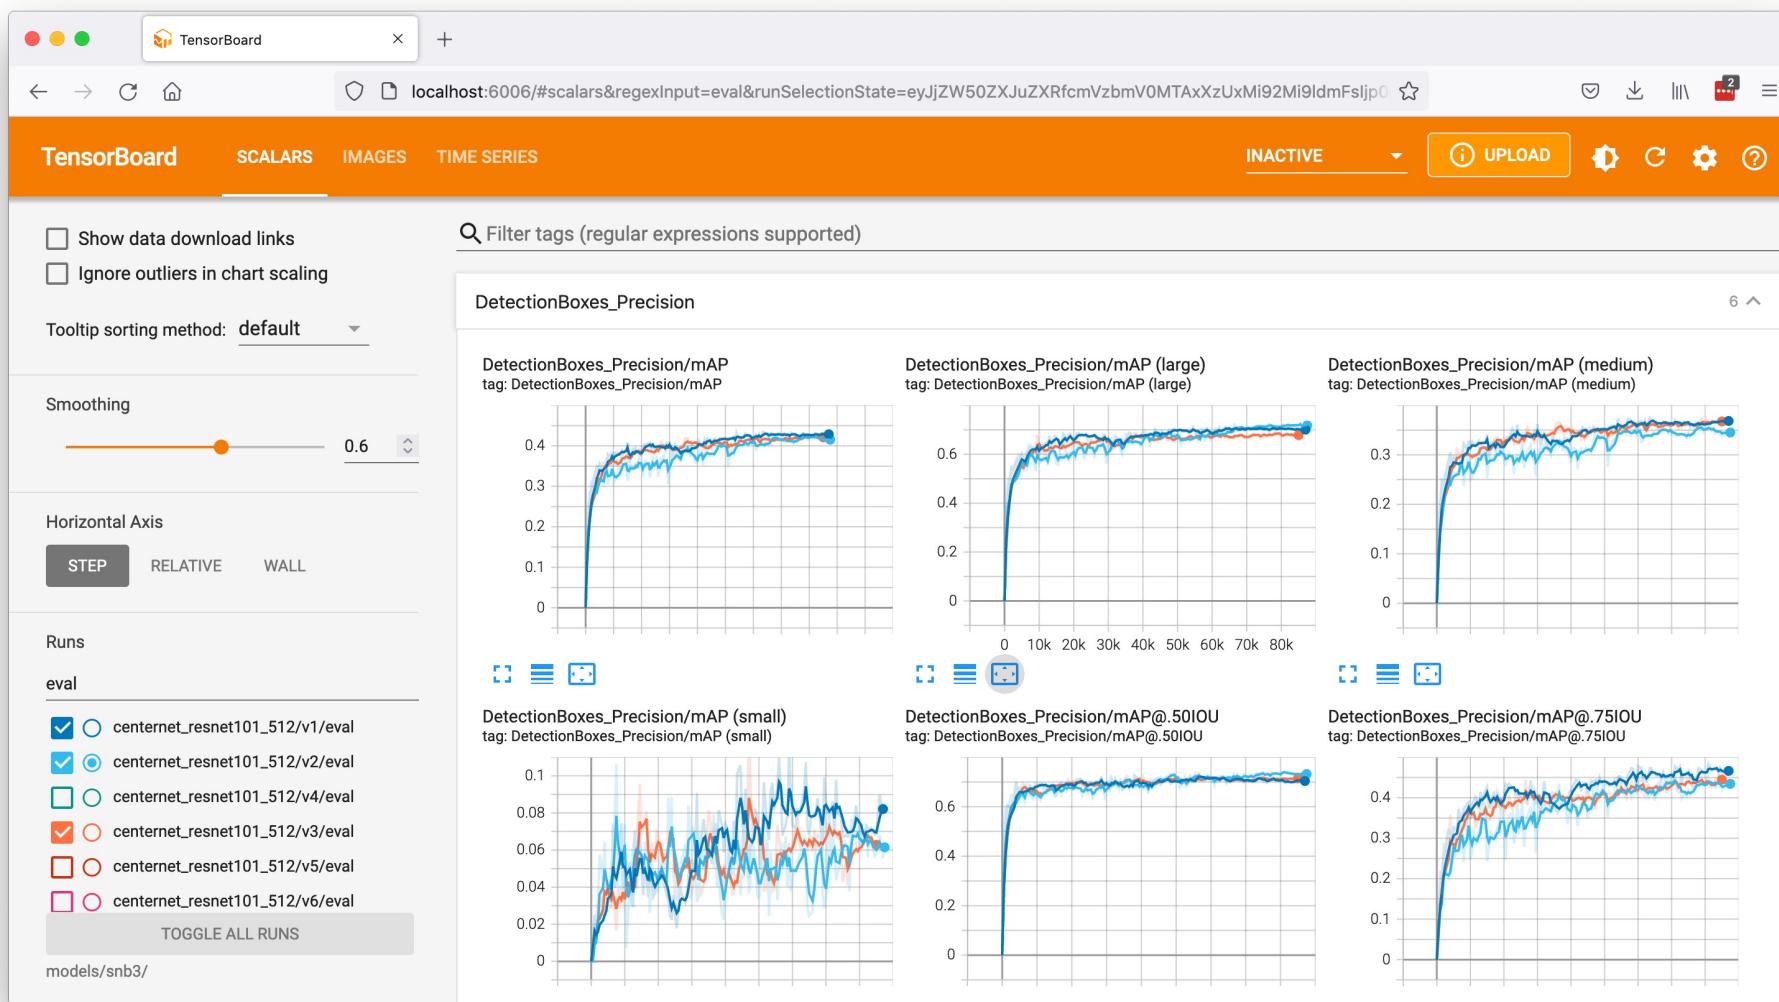
# EXPLORING TENSORBOARD

# EXPLORING TENSORBOARD

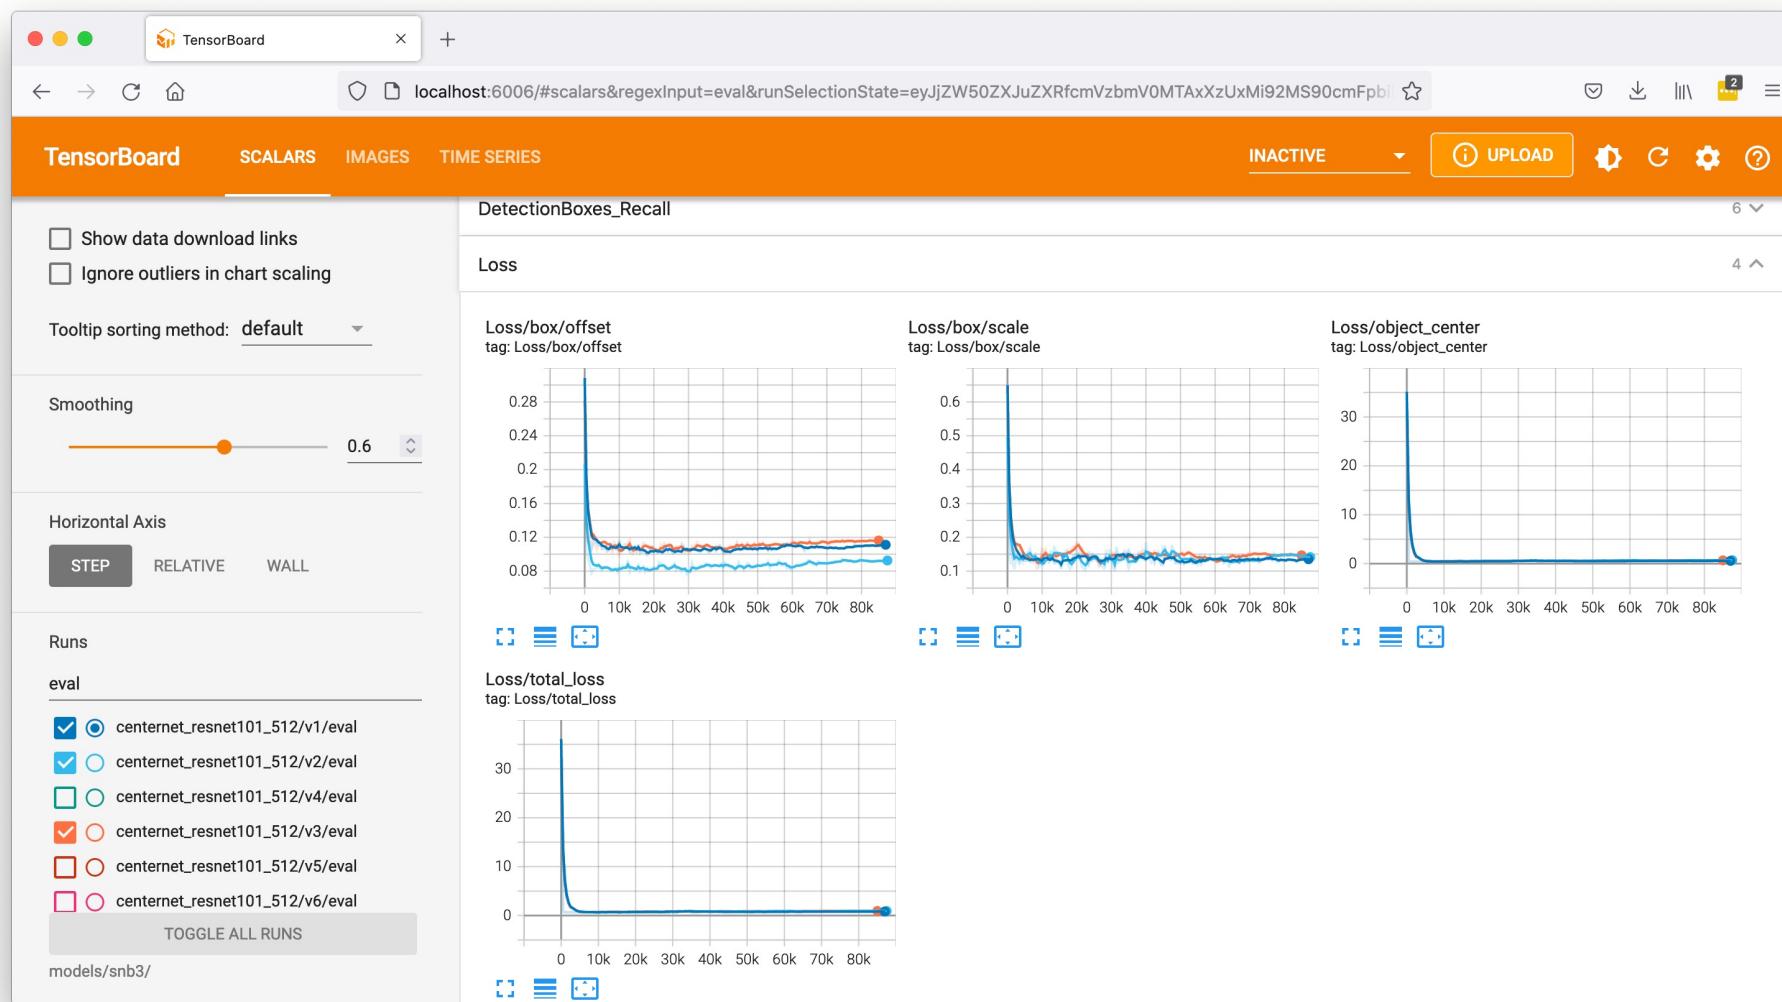
- Scalars
- Images
- Time Series



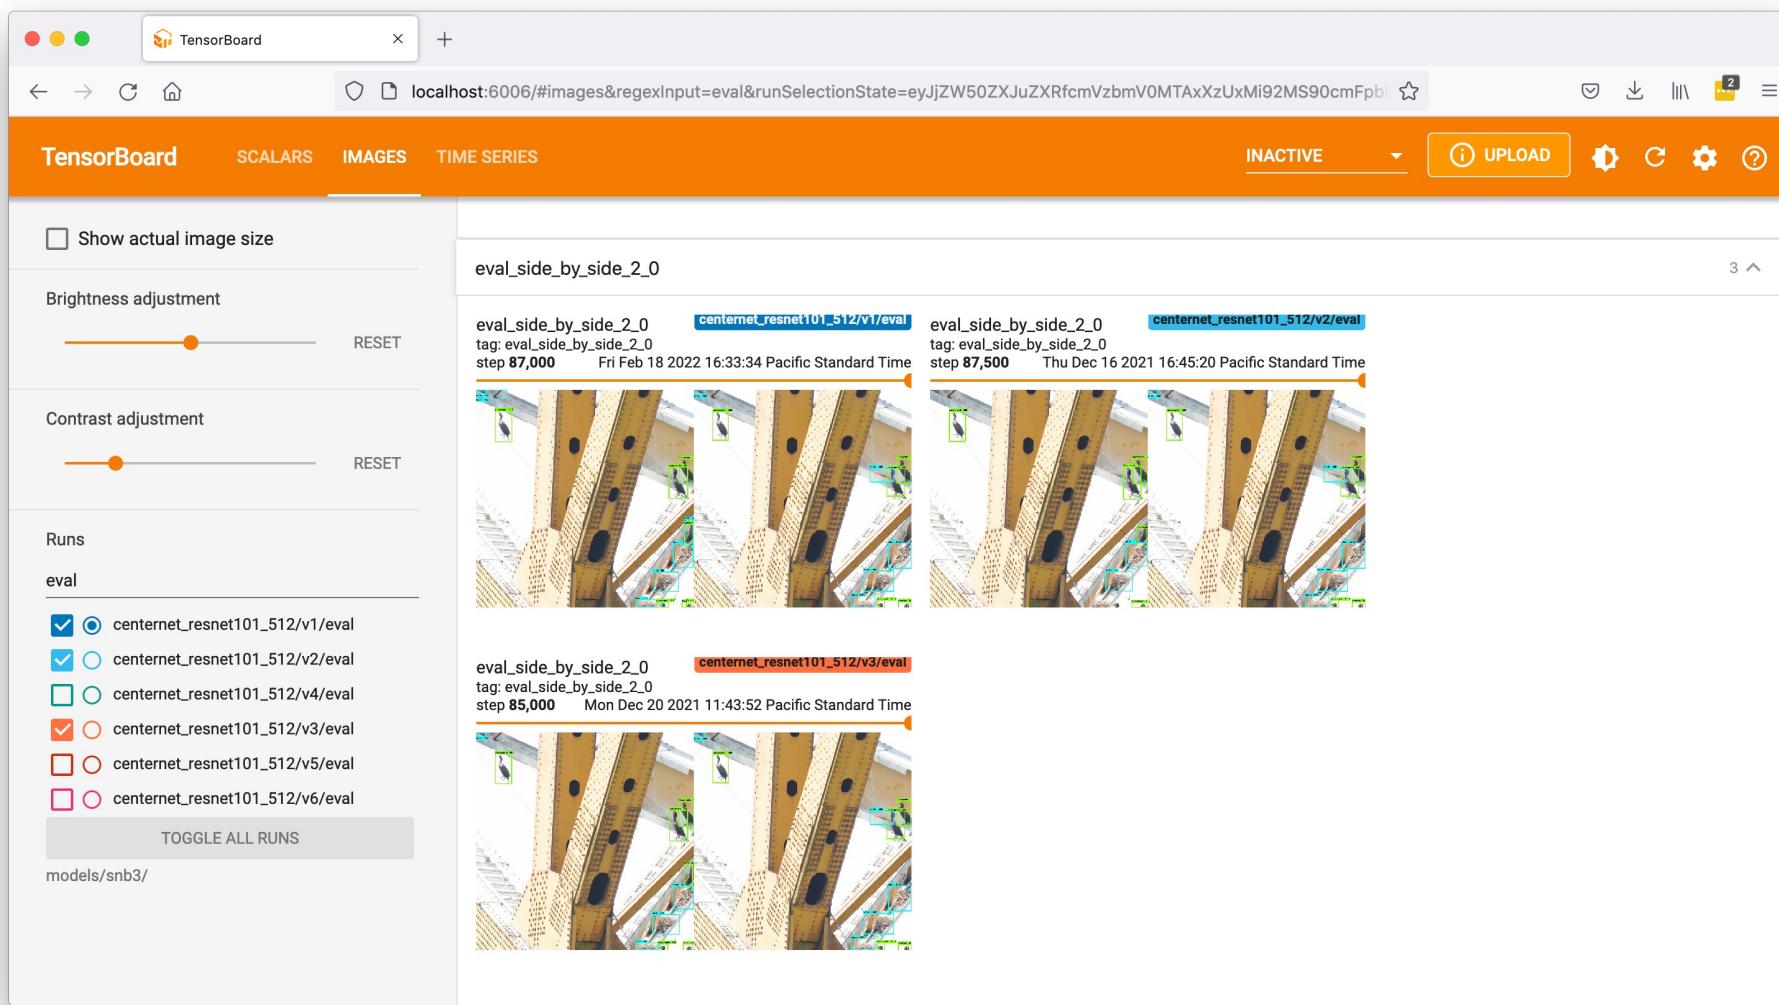
# SCALARS



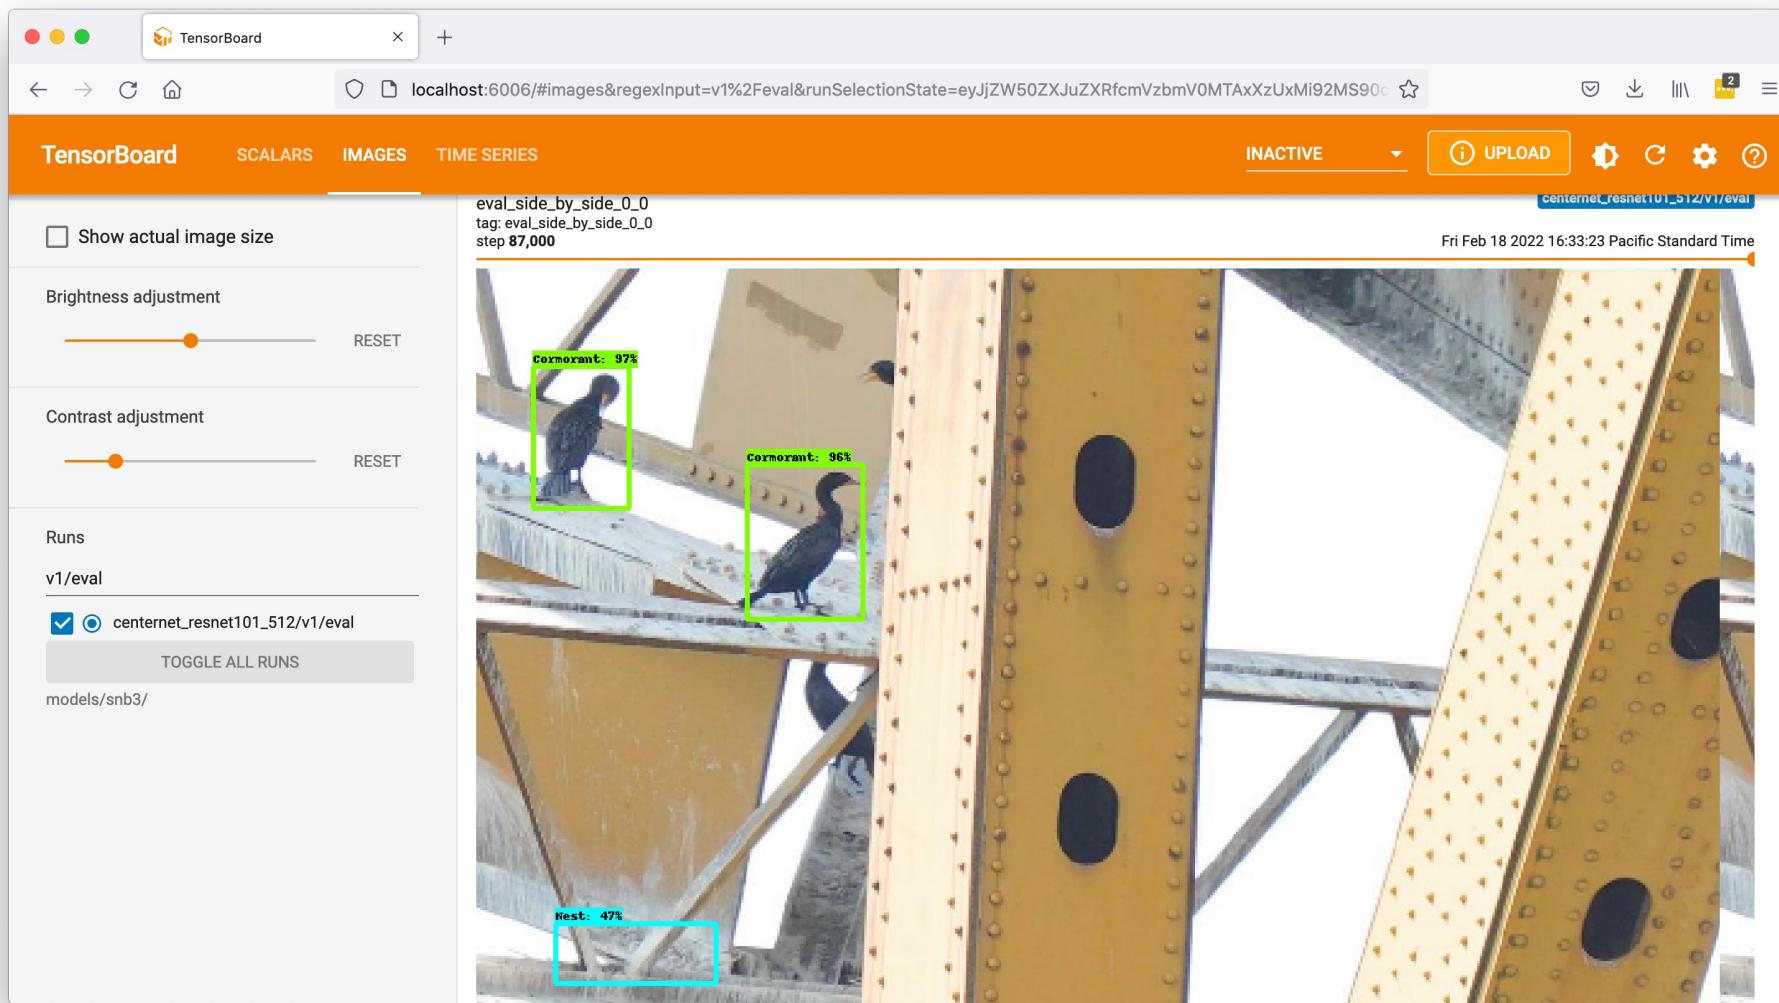
# SCALARS



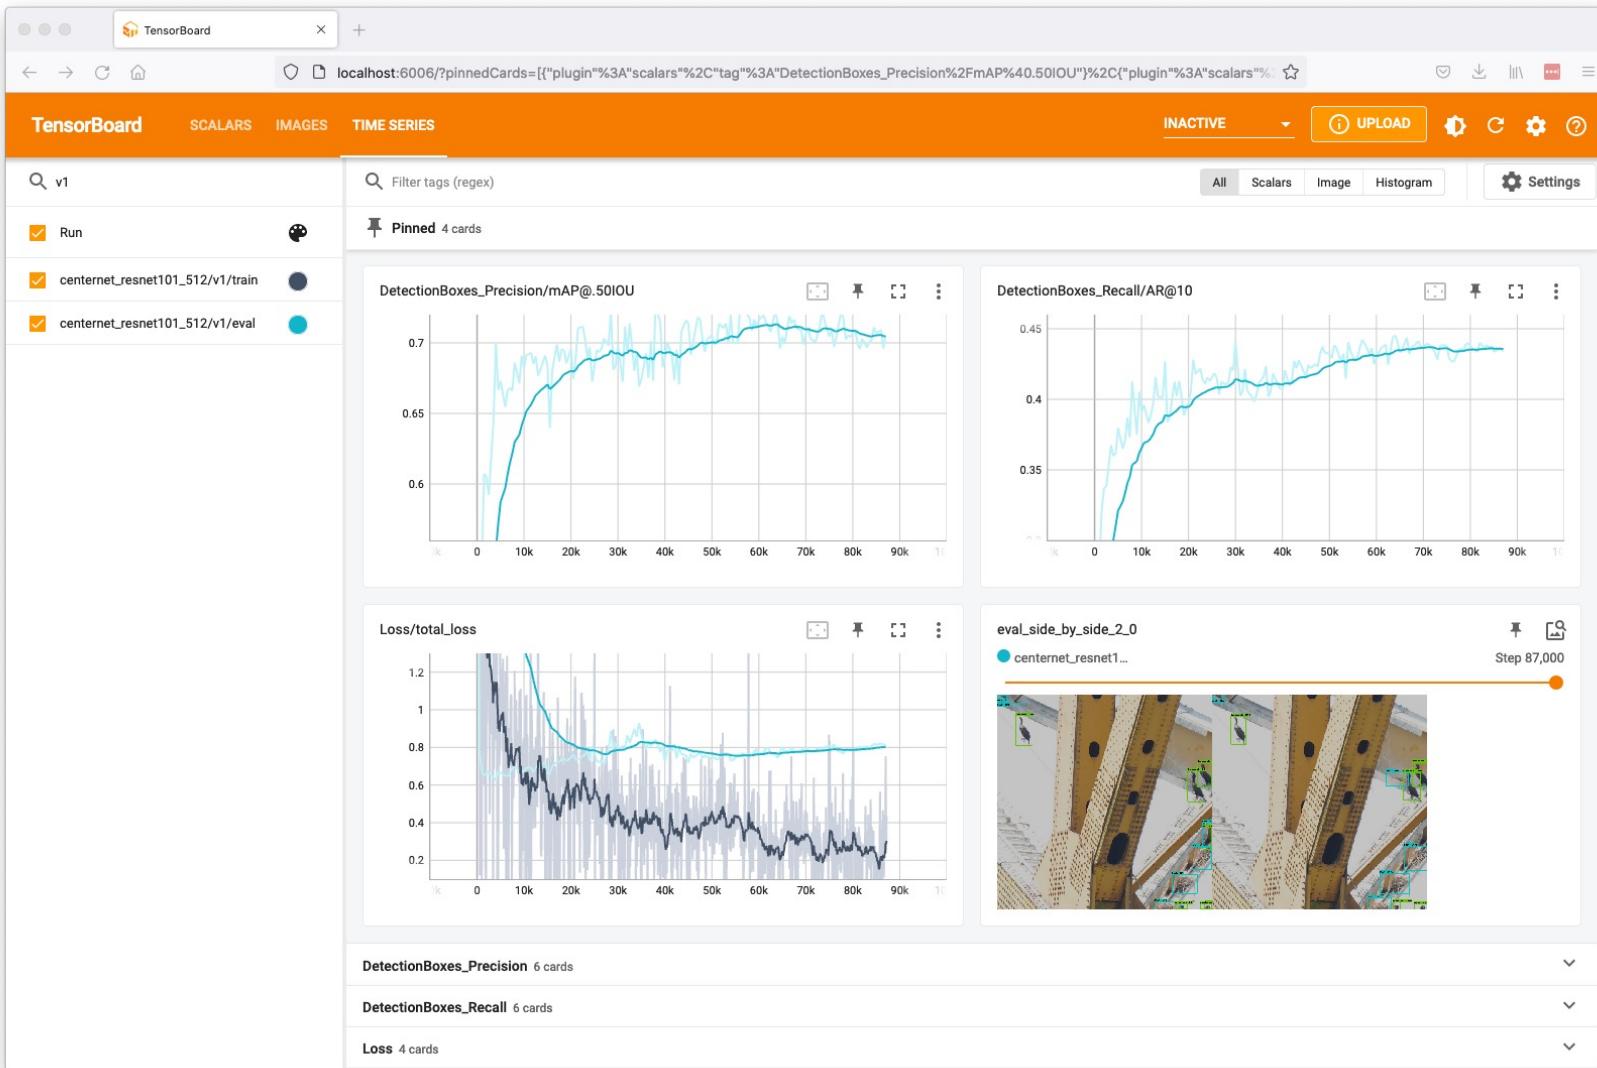
# IMAGES



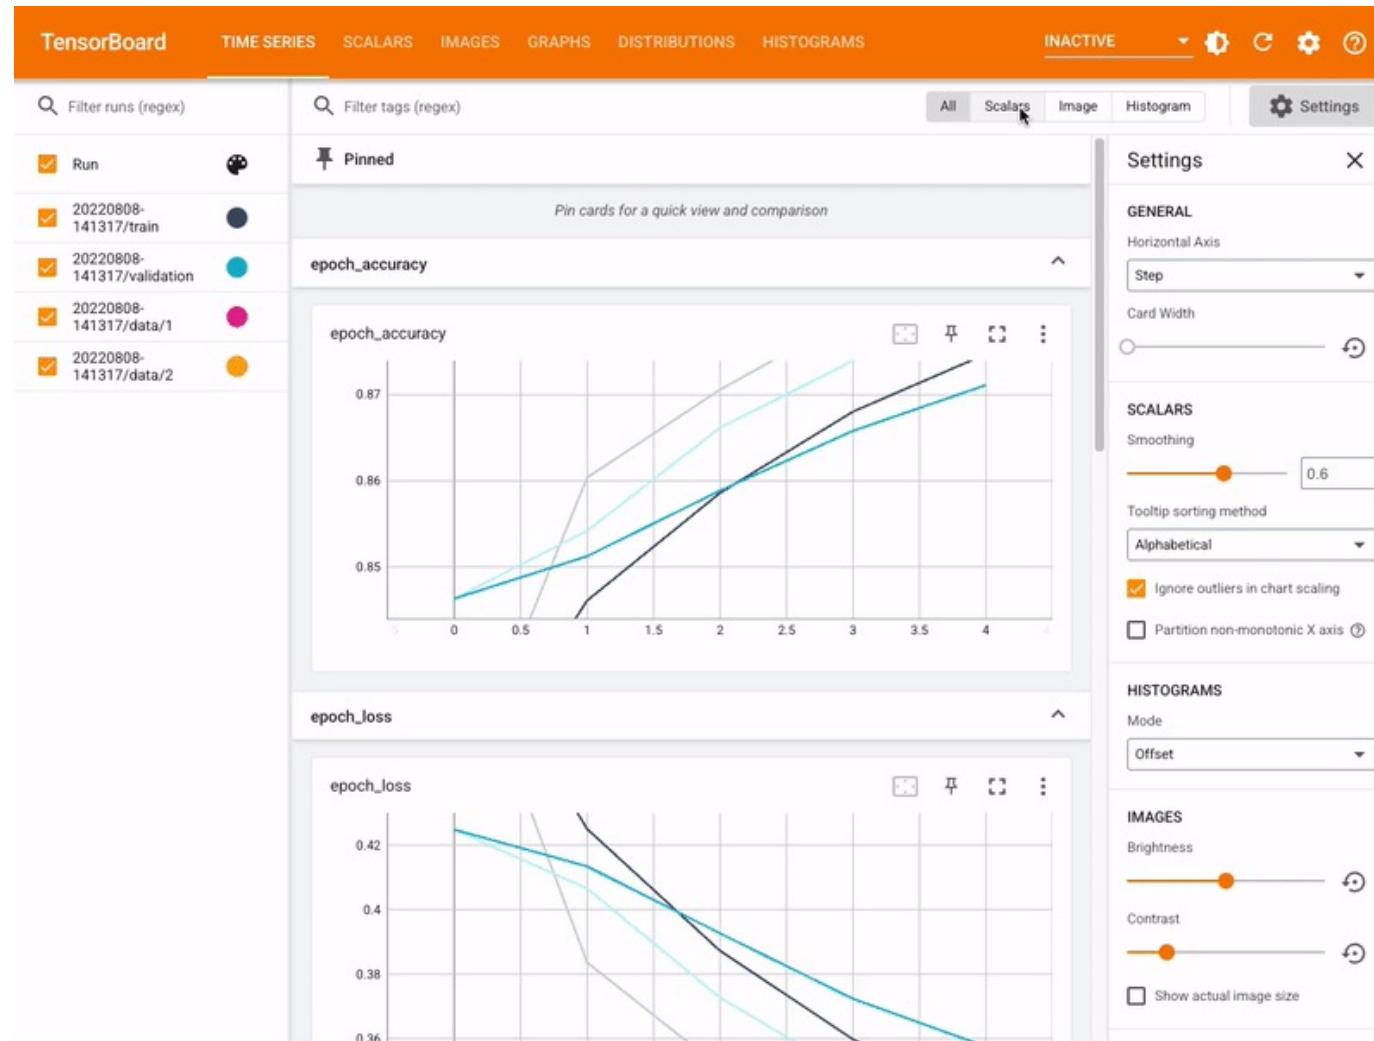
# IMAGES



# TIME SERIES



# OTHER TENSORBOARD DASHBOARDS



**THANK YOU**

**QUESTIONS?**