

- Current maintenance

- ▶ Cedar is down for system software updates (today only)
- ▶ `/project` expansion March 1<sup>st</sup> - 4<sup>th</sup>, Lustre file metadata will be copied over to new SSDs, `/project` unavailable at this time, you can still use `/scratch` for running jobs during this time

- Cedar is a very complex system

- ▶ ~ 1,600 nodes on Cedar, each with 24-48 cores, local storage
- ▶ ~ 66,000 cores
- ▶ two different CPU architectures
- ▶ 100Gb/s Omnipath interconnect linking all nodes and storage
- ▶ three (Lustre) parallel file systems with ~ 14 PB combined storage
  - `/home`, `/scratch`, `/project`
  - each with its own policies, 2/3 backed up
- ▶ 584 NVIDIA P100 Pascal GPUs
- ▶ ~ 60 Slurm partitions, for long / short / GPU / large-memory / interactive jobs / CPU architecture
  - 28d, 7d, 3d, 1d, 12h, 3h maximum runtime
  - trying to accommodate a large variety of job types
  - at the cost of efficiency and simplicity

- Our goals are to:

- ▶ provide as much uptime as possible
  - we constantly monitor our clusters
  - work as quickly as possible to repair problems and return nodes to production
  - in case of downtime or other problems, provide frequent system status updates
- ▶ accommodate a wide spectrum of jobs
- ▶ maximize resource (CPUs, GPUs, memory, storage) utilization
- ▶ minimize turnaround for your jobs

① When hardware/etc problems occur, we want you to know how *in some cases* you can work around them

② We want to show you how certain workflows can lead to problems on HPC clusters


- ▶ and share with you best practices for working on these systems

## Node failures: a node needs rebooting or other work

- ~ 1,600 nodes on Cedar
- Of these, ~ 30 nodes have actual hardware failures at any one time
  - ▶ these get gradually replaced through a rather onerous return merchandise authorization (RMA) process
  - ▶ we are working with the vendor to simplify this process
- Marked offline by Slurm, for any number of reasons: not communicating, incorrect reports, low memory, cannot terminate the job, etc.
  - ▶ requires manual intervention
- Over-subscription of nodes, GPUs
  - ▶ e.g., too many threads
- Does not pass other checks and taken offline
  - ▶ GPUs get stuck in a strange state: *"Only EGL 1.4 and greater allows OpenGL as client API"*, requires reboot

## File system problems

- Lustre object storage servers (OSS) can get overloaded with lots of small requests
  - ▶ example: this past Tuesday a user was running 90 jobs, all with high I/O in `/project` bringing it to a halt
  - ▶ putting these jobs on hold *did not* fix the system
  - ▶ OSS had to be rebooted
  - ▶ end result: `/project` was not available to all users for a couple of hours
- On Cedar we have:
  - ▶ 4 object storage servers handling `/home` (slow) and `/scratch` (fast)
  - ▶ 10 object storage servers handling `/project`
- These are paired into groups of two
  - ▶ one in a pair goes down  $\Rightarrow$  the other one will take over, but high I/O jobs might take much longer than expected
  - ▶ both go down  $\Rightarrow$  the entire filesystem will hang

more on high IOps later 

## Scheduler failures

- Slurm can get overloaded with too many requests

more on scheduler later 

- No software stack synchronization between login and compute nodes
- Networking problems (within or outside our control)


# What do you see?


- Sluggish jobs
- Jobs not starting / taking unusually long to start
  - ▶ also valid reasons why your job's estimated start time could be pushed into the future
- Slurm not responding, or producing unusual output
  - ▶ e.g. last year's infamous Slurm bug leading to jobs stuck in 'Prolog' R (running) state for a long time, not producing any output
- Shell not responding to simple commands or very slow
  - ▶ could be per individual filesystem/command
- Output files missing from your working directory
- Inside running jobs see "*module not found*"
  - ▶ typically requires manual intervention
- Cannot log in

# What can you do about these instabilities?



- Report problems to `support@computecanada.ca` with details:
  - ▶ system you are using
  - ▶ job IDs of affected jobs
  - ▶ detailed description of the problem
  - ▶ time/date it was first encountered
  - ▶ full path to one of the directories with the script and error files
    - check if you signed the consent that allow analysts to check your files (this will help resolve problems quickly instead of exchanging many emails), by logging in to `http://ccdb.computecanada.ca` and selecting My Account ↗  
Agreements
- Pay attention to login messages
  - ▶ terminal output from anything in your `~/.bash_profile` or `~/.bashrc` (e.g. when loading a module or activating a virtual environment ) might force important system messages scroll past the top of the terminal
  - ▶ these may contain both general system notices and `/scratch` purge notifications specifically for you
- Check `http://status.computecanada.ca` for updates and recent incidents

Yes, I allow Compute Canada team members to access my files on Compute Canada systems as part of an on-going support request as described above. 

No, please ask me every time. 

# What can you do? (cont.)

- Sometimes you could work around a temporary filesystem problem by submitting jobs from another filesystem
  - ▶ on Cedar `/home,/scratch` files are handled by different servers than `/project` (may not be always possible: performance, input data)
- Do *not* delete and resubmit jobs that have been waiting in a queue for a long time until confirming with `support@computecanada.ca`
  - ▶ otherwise we can't analyze why a job is waiting, and priority may be lost
- Expect a backlog of jobs after a system problem
  - ▶ do *not* swamp the system with a bunch of new jobs – be selective about what is most important to you
  - ▶ make sure that job parameters are chosen carefully to match the needs of particular jobs



# These workflows will create problems

- Running anything CPU-intensive on the head node
- Submitting large number of jobs
- Issuing too many requests to the scheduler
  - ▶ classical example: running `watch squeue ...` (never do this!)
  - ▶ submitting thousands of jobs and then cancelling them
- Complex/unrealistic job dependencies can make Slurm unstable
- Not testing first on a small scale, and not scaling up gradually
  - ▶ large parallel jobs
  - ▶ many serial jobs and large job arrays
  - ▶ large computational problems in general
- Assuming perfect parallel scaling
  - ▶ your 64-core job may be slower than 32-core ...

- Excessive and/or “bad” I/O, i.e. anything resulting in high load on Lustre object storage servers
  - ▶ avoid high I/O workflows: lots of small files, read/write in chunks smaller than 1MB, reading small blocks from large files
- Storing a large number of small files
  - ▶ organize your code’s output
  - ▶ use **tar**, or even better **dar** (<http://dar.linux.free.fr>, supports indexing, differential archives, encryption)
- Using nested parallelism in black-box pipelines
  - ▶ e.g. submitting serial jobs each of which launches multiple threads, sometimes asking for all cores on a node
  - ▶ your pipeline should be adapted to the cluster; if not sure, please talk to us


- Using `mv` command to move files `/home,/scratch` → `/project` will result in an overquota error message in the middle of moving
  - ▶ this is expected behaviour!
  - ▶ not so much a problem for the cluster, but certainly will be a problem for you ...
  - ▶ in `/project` the 1TB (or higher) quota is applied to all files with the group ID `def-group`
    - so that all your group members are able to write there
    - any new file you write to `/project` will have `def-group` group ID
    - you can find this group ID by running `id` and looking for 'def-...'
  - ▶ by default, all files in `/home,/scratch` have group ID `username`
  - ▶ `mv` command preserves group ID, i.e. effectively `mv` acts as `cp -a`
  - ▶ the quota for group ID `username` is almost zero in `/scratch`
  - ▶ solution: use `cp` instead, followed by `rm`

- Implement/use checkpointing to be prepared for system failures
- Break your job into pieces, if possible (time-wise, processor-wise)
- Read the documentation about scheduling, running jobs, using modules, other topics [📖 https://docs.computecanada.ca](https://docs.computecanada.ca)
- Know as much as possible about your application (serial vs. parallel), and how it was parallelized (threaded vs. MPI)
  - ▶ very important for creating the correct job submission script!
- Start with some tests before running extensive simulations
  - ▶ estimate the resources (especially memory, wall time)
  - ▶ use `sacct` or `seff` to estimate your completed code's memory usage
  - ▶ test parallel scaling, scaling with problem size
- Only request resources (memory, running time) needed
  - ▶ with a bit of a cushion, maybe 115-120% of the measured values

# Other best practices (cont.)

- If you still need to do lots of small I/O from inside your job:
  - ▶ use Slurm-generated directory `$SLURM_TMPDIR` (pointing to `/localscratch/${USER}.${SLURM_JOBID}.0` on a node's SSD) for both input and output
    - don't forget to move files out before your job terminates: everything in `$SLURM_TMPDIR` will be deleted
  - ▶ use `$TMPDIR` RAM disk (pointing to `/tmp`)
    - don't forget to allocate additional memory to your job
    - don't forget to move the results before your job terminates
- Port your workflow to another CC's general-purpose cluster, to run it there in case of failures
  - ▶ data management part may not be so easy, but Globus should help
  - ▶ also try to port your workflows (have accounts, appropriate input data, programs installed) to local clusters where available (Grex, Orcinus, Plato)
- If you received a `/scratch` purge warning, do *not* wait until the last minute to transfer data to local systems or other clusters
  - ▶ always pay attention to `/scratch` purge notices (email and system login message)
  - ▶ exercise care when transferring data close to quota

# Other best practices (cont.)

- Be aware that some filesystems are not backed up (e.g. `/scratch`), and some have a purge policy (`/scratch`)  **have a backup plan**
- If a file's path changes, our backup system will interpret it as a new file  
⇒ unnecessary load on the filesystems
  - ▶ be careful with renaming large directories in `/home` and `/project`
- In general, do *not* run jobs in `/home`
  - ▶ slow, not designed for high performance (unlike `/scratch`)
  - ▶ small quota (50GB/user)
  - ▶ lots of I/O makes difficult to do backups
- After your job finishes:
  - ▶ clean up (remove files that are no longer needed)
  - ▶ compress large files to reduce the disk space usage
  - ▶ archive (tar) the directories with many small files to reduce the file count
  - ▶ eventually move your data from `/scratch` to `/project`, `~/nearline` (will be available on Cedar soon), your own storage